

Forward KL Can Be Better: Normalizing Flow Boltzmann Sampling with Log-Ratio Variation

Qi Feng¹, Rongjie Lai², Di Qi², and Xuda Ye²

¹Department of Mathematics, Florida State University, Tallahassee, FL 32306


²Department of Mathematics, Purdue University, West Lafayette, IN 47907

July 22, 2026

Abstract

Normalizing flow Boltzmann generators retain a tractable proposal density, but training with forward KL depends on target samples that may be biased or omit modes. A flow can then miss target mass while its observed importance weights give a high effective sample size. We introduce log-ratio variation, a pairwise regularizer of the exact target-to-model log-density ratio, and use it to define KLXX. Its three weighting measures provide signals for target fitting, candidate mode coverage, and leakage control without changing the exact-target minimizer. We derive an exponential Fisher–Rao bound and a fixed-surrogate error bound for forward KL with the target-weighted term, then place KLXX in an adaptive-staging Boltzmann generator with exact importance reweighting and sharpening of the molecular potential. In the reported model and molecular tests, KLXX improves coverage-aware or stagewise diagnostics, and selected molecular observables agree with independent OpenMM simulations. Log-ratio variation thus supplies information that a sampled forward KL objective can omit.

KLXX: one log-ratio (KL), two log-ratio variations (XX)



1 Introduction

Drawing samples from a target distribution $\mu \propto \exp(-U)$ on $\mathbb{R}^d$ is a recurring problem in molecular simulation, Bayesian inference, and statistical physics (Frenkel and Smit, 2002; Robert and Casella, 2004; Landau and Binder, 2014). When μ is high-dimensional and multimodal, its mass splits among basins of U separated by barriers. Local sampling kernels, including Metropolis–Hastings (Metropolis et al., 1953; Hastings, 1970), Langevin dynamics (Langevin, 1908; Leimkuhler and Matthews, 2015), Nosé–Hoover dynamics (Nosé, 1984; Hoover, 1985), and Hamiltonian Monte Carlo (Duane et al., 1987), evolve from the current configuration and can therefore remain confined to one basin over accessible simulation times. When this occurs, finite-time estimates are biased toward the visited basin, with no direct account of the mass not reached.

Tempering and particle methods address this problem by connecting a tractable distribution to μ through intermediate distributions. Parallel tempering (PT) and replica exchange run copies of the system at several temperatures and exchange neighboring configurations, so a basin reached at high temperature can enter the low-temperature ensemble (Swendsen and Wang, 1986; Geyer, 1991; Hukushima and Nemoto, 1996). Annealed importance sampling (AIS) moves along a bridge of intermediate distributions while accumulating importance weights (Neal, 2001). Sequential Monte Carlo (SMC) uses a similar bridge but resamples and rejuvenates a weighted particle set at each step (Doucet et al., 2001; Del Moral et al., 2006). For a specified chemical process, umbrella sampling biases selected reaction coordinates to cross relevant barriers and reweights the biased samples to recover a free-energy profile (Torrie and Valleau, 1977; Kästner, 2011). Nested sampling instead rewrites the normalizing integral as a one-dimensional quadrature over a likelihood threshold (Skilling, 2006).

Recent work improves PT through three distinct mechanisms. Nonreversible swap schedules and tuned temperature grids accelerate round trips (Syed et al., 2022); variational-reference methods adapt the reference endpoint and its connecting path (Surjanovic et al., 2022); and neural transports improve communication between adjacent distributions with limited overlap (Zhang et al., 2025). All three strengthen communication between neighboring distributions, but none removes the dependence on local exploration: a basin not reached by any chain cannot be passed to the target ensemble (Woodard et al., 2009; Chehab et al., 2025). Improving this communication therefore does not by itself resolve missing-mode coverage. This gap motivates a learned proposal with a tractable density whose training draws information from more than one sampling measure.

Boltzmann generators provide a complementary route through a learned transport map (Noé et al., 2019). Let $\mu_0 \propto \exp(-U_0)$ be an easily sampled reference distribution, usually Gaussian, and let $G : \mathbb{R}^d \rightarrow \mathbb{R}^d$ be an invertible normalizing flow (Rezende and Mohamed, 2015; Papamakarios et al., 2017; Kingma and Dhariwal, 2018). We use the target-to-source convention $G_{\#}\mu \approx \mu_0$. The inverse map then sends independent source draws toward the target, and their pushforward distribution is $\nu := (G^{-1})_{\#}\mu_0$. Its density function is

$$\nu(y) = \mu_0(G(y)) |\det J_G(y)|, \quad (1)$$

where $J_G(y)$ is the Jacobian matrix of G at y . An importance weight proportional to μ/ν is therefore available for every pushforward sample without fitting another density or tracing a stochastic path. These weights enable importance-reweighted estimates and define the effective sample size (ESS), a diagnostic of density agreement on the support reached by ν (Liu, 2001; Doucet et al., 2001; Elvira et al., 2022). The remaining challenge is that a high ESS certifies agreement only on the reached support: a target mode absent from the pushforward samples contributes no observed weight.

The choice of training objective determines how this challenge enters training. The original Boltzmann generator (Noé et al., 2019) and its variational-inference predecessor (Rezende and Mohamed, 2015) minimize the reverse Kullback–Leibler (KL) divergence

$$\text{(reverse)} \quad \text{KL}(\nu \parallel \mu) = \int_{\mathbb{R}^d} \nu(x) \log \frac{\nu(x)}{\mu(x)} dx, \quad (2)$$

which needs samples only from ν and evaluations of U . Its empirical gradient is therefore accessible, but the divergence is mode-seeking (Wainwright and Jordan, 2008): a mode absent from ν contributes no samples and hence no gradient. A flow can then collapse onto a subset of the target modes (Wu et al., 2020; Felardos et al., 2023). Existing approaches interleave MCMC kernels with deterministic layers, encode target symmetries in

the architecture, reduce gradient variance near the optimum, or anneal from a smoother potential (Wu et al., 2020; Köhler et al., 2020; Vaitl et al., 2022; Schopmans and Friederich, 2025). These changes can improve exploration, but their success still depends on whether the resulting sampler reaches each relevant mode.

The opposing objective minimizes the forward KL,

$$\text{(forward)} \quad \text{KL}(\mu\|\nu) = \int_{\mathbb{R}^d} \mu(x) \log \frac{\mu(x)}{\nu(x)} dx, \quad (3)$$

which is mass-covering for exact distributions: it diverges if ν vanishes where μ has positive mass (Wainwright and Jordan, 2008). Its empirical form, however, needs samples from μ , the distribution being learned. Existing work resolves this circularity in several ways. Maximum likelihood uses configurations from a reference thermodynamic state (Wirnsberger et al., 2020); adaptive flow samplers train on a running Markov chain and use the flow as a global proposal (Gabri   et al., 2022); molecular flows can fit a molecular-dynamics trajectory (Klein et al., 2023); and transferable architectures amortize coverage across related molecules (Klein and No  , 2024). Persistent chains and coupled-target regression provide other approximations to the missing target expectation (Naesseth et al., 2020; Rehman et al., 2026). Because every empirical loss can see only the regions represented by its sample source, this exact-distribution property does not prevent empirical mode omission. This tension motivates an objective that retains forward KL while obtaining training signals beyond one target surrogate.

The sampled forward KL leaves three distinct gaps. First, a mode absent from its target surrogate contributes no training points. If the flow also misses that mode, the observed weights can be nearly uniform on the reached support, so the ESS can be close to one even though target mass is absent; we call this a misleading or *fake ESS* (No   et al., 2019; Midgley et al., 2022). Second, a finite or biased target surrogate distorts the sampled objective, and the fitted flow can inherit that error. Third, target samples give little information where the proposal places excess mass, allowing leakage between modes or into target-poor regions. Improving one of these failures does not by itself resolve the other two.

To address these gaps, we propose KLXX, a forward KL loss augmented by log-ratio variation; the graphical abstract summarizes it, and Equation (8) gives its full default form. KLXX preserves the tractable target-to-proposal density ratio and exact-target minimizer while separating target fitting, candidate-mode coverage, and proposal-support control. Annealing supplies the target surrogate and target-weighted variation. The new third term mixes quench and temper (QT) and pushforward samples: QT contributes candidate modes missing from the surrogate, while pushforward samples expose proposal support and leakage. QT guides coverage rather than estimating physical mode weights or guaranteeing complete discovery.

The closest objective-level comparison is Flow Annealed Importance Sampling Bootstrap (FAB), which couples AIS and flow training but replaces the forward KL by the α -divergence ($\alpha = 2$) (Midgley et al., 2022). Its integrand is proportional to μ^2/ν , and its exact minimizer also minimizes importance-weight variance. FAB targets μ^2/ν with AIS and reuses those samples through a prioritized replay buffer, adding buffer maintenance. Related methods insert flows into SMC, anneal toward low-temperature targets, add an inference-time SMC correction, or connect disconnected regions (Arbel et al., 2021; Matthews et al., 2022; Schopmans and Friederich, 2025; Tan et al., 2025; von Klitzing et al., 2026). KLXX instead retains the forward KL and adds nonnegative log-ratio variations. Its extra cost is QT and the variation evaluations, including the additional mixture batch; it uses no replay buffer. FAB’s 2-divergence gives a stronger tail penalty and may reduce the need for separate coverage construction, whereas KLXX obtains coverage information from QT. Remark 3 compares the sampling mechanisms after they are defined.

Centered log-ratio penalties appear in log-variance objectives for learned diffusion samplers (Richter and Berner, 2024) and in LDR for data-based Boltzmann generator objectives (Schopmans et al., 2026); such centered log-ratio regularization in Boltzmann generator training is therefore not new. Developed independently of LDR, our target-weighted term instead pairs the forward KL integrand at two independent target points, yielding the double integral below. The new KLXX component is the variation term weighted by a mixture of QT and pushforward samples.

A separate line trains a stochastic process rather than a bijection. Controlled stochastic differential equations, denoising diffusions, and flow-matching models learn a drift or velocity that transports the reference toward μ (Zhang and Chen, 2022; Vargas et al., 2023; Berner et al., 2024; N  sken and Richter, 2021; Richter and Berner, 2024). Energy-only variants such as iterated denoising energy matching (iDEM) alternate a sampling loop with a score- or velocity-fitting loop and a replay buffer (Akhound-Sadegh et al., 2024; Woo

and Ahn, 2024). Learned nonequilibrium transport augments AIS with a trainable drift, and later work extends this approach to larger many-body systems (Albergo and Vanden-Eijnden, 2025; Havens et al., 2025; Akhound-Sadegh et al., 2025).

These stochastic transports can be more expressive than the invertible architectures used here. Their terminal density, however, often requires per-sample ODE integration (He et al., 2025). We retain a bijection because the exact proposal density, importance ratio, ESS, and reweighting are central to the objective and final estimator. Table 1 summarizes these distinctions.

method	objective	training mechanism
reverse KL	$\text{KL}(\nu\ \mu)$	pushforward samples; no target sampler
forward KL	$\text{KL}(\mu\ \nu)$	annealed target surrogate in this work
FAB	2-divergence (μ^2/ν)	AIS on μ^2/ν ; MCMC differentiates ν ; replay buffer
LDR	$\text{KL}(\mu\ \nu) + \mathbb{E}_\mu[z - \mathbb{E}_\mu z ^p]$	fixed energy-labeled data; no extra proposal samples
iDEM	energy matching	iterative sampler; stochastic transport; replay buffer
KLXX	$\text{KL}(\mu\ \nu) + X_\mu + X_{(\hat{\mu} + \bar{\nu})/2}$	target annealing, QT, and detached pushforward samples no MCMC derivative of ν or replay buffer

Table 1: Objectives and training mechanisms for the methods discussed above. Here $z = \log(\mu/\nu)$; the LDR row gives its target-weighted form with $p = 1$ for LDR-1 and $p = 2$ for LDR-2. The forward KL row uses the annealed baseline in this paper but may instead use external target data. The table compares how samples enter each objective; it does not assume equal computational budgets.

We now turn from the principle-level comparison to the mathematical definition of KLXX. Define the tractable log-density ratio $z(x) := \log \frac{\mu(x)}{\nu(x)}$. For a nonnegative density ω on $\mathbb{R}^d$, define the *log-ratio variation*

$$X_\omega(\mu\|\nu) = \iint_{\mathbb{R}^d \times \mathbb{R}^d} \omega(x) \omega(y) \left| \log \frac{\mu(x)}{\nu(x)} - \log \frac{\mu(y)}{\nu(y)} \right| dx dy, \quad (4)$$

which measures how far z is from a constant on the region weighted by ω . For $\omega = \mu$, X_μ averages $|z(x) - z(y)|$ over independent target points, whereas LDR-1 uses the centered penalty $\mathbb{E}_\mu |z - \mathbb{E}_\mu z|$. Jensen’s inequality and the triangle inequality give $\mathbb{E}_\mu |z - \mathbb{E}_\mu z| \leq X_\mu(\mu\|\nu) \leq 2\mathbb{E}_\mu |z - \mathbb{E}_\mu z|$. Thus, at the same coefficient, KL + X_μ and LDR-1 are comparable up to a factor of two. This suggests similar training performance, although the objectives are not equivalent. KLXX is also not equivalent to LDR-2: its target-weighted variation is pairwise L^1 , and its new mixture term uses QT and pushforward samples. Since X_ω is nonnegative and vanishes at $\nu = \mu$, adding it to the exact forward KL preserves the target as a minimizer.

The target distribution μ supplies samples for both the forward KL and X_μ . The wide-coverage distribution $\hat{\mu}$ supplies candidate-mode samples, while $\bar{\nu}$ is the pushforward distribution of the current proposal; samples from $\bar{\nu}$ are held fixed during differentiation. The combined loss is

$$\mathcal{L}[\nu] = \text{KL}(\mu\|\nu) + \lambda X_\mu(\mu\|\nu) + X_{\alpha\hat{\mu} + \beta\bar{\nu}}(\mu\|\nu), \quad (5)$$

where $\lambda, \alpha, \beta \geq 0$. For simplicity, the introduction assumes $\alpha + \beta = 1$; Section 3.1 allows arbitrary non-negative mixture weights. Under this convention, the final term is the log-ratio variation under the mixture distribution $\alpha\hat{\mu} + \beta\bar{\nu}$. Its pairwise expectation includes pairs drawn within either distribution and pairs with one sample from each distribution.

Although training seeks $G_\# \mu \approx \mu_0$, the pushforward definition implies the exact identity $G_\# \nu = \mu_0$. Combining (1) with $\mu(y) = \exp(-U(y))$ and $\mu_0(x) = \exp(-U_0(x))$ yields

$$z(y) := \log \frac{\mu(y)}{\nu(y)} = U_0(G(y)) - U(y) - \log |\det J_G(y)|. \quad (6)$$

Hence the loss in terms of the map is

$$\mathcal{L}[G] = \underbrace{\mathbb{E}_{y \sim \mu} [z(y)]}_{\text{KL}(\mu\|\nu)} + \lambda \underbrace{\mathbb{E}_{y, y' \sim \mu} [|z(y) - z(y')|]}_{X_\mu(\mu\|\nu)} + \underbrace{\mathbb{E}_{y, y' \sim \alpha\hat{\mu} + \beta\bar{\nu}} [|z(y) - z(y')|]}_{X_{\alpha\hat{\mu} + \beta\bar{\nu}}(\mu\|\nu)}. \quad (7)$$

The default choice $(\lambda, \alpha, \beta) = (1, \frac{1}{2}, \frac{1}{2})$ is

$$\text{KLXX}[G] = \underbrace{\mathbb{E}_{y \sim \mu} [z(y)]}_{\text{KL}(\mu \parallel \nu)} + \underbrace{\mathbb{E}_{y, y' \sim \mu} [|z(y) - z(y')|]}_{X_\mu(\mu \parallel \nu)} + \underbrace{\mathbb{E}_{y, y' \sim (\hat{\mu} + \bar{\nu})/2} [|z(y) - z(y')|]}_{X_{(\hat{\mu} + \bar{\nu})/2}(\mu \parallel \nu)}. \quad (8)$$

With equal target and mixture batch sizes, KLXX evaluates about twice as many points per step as forward KL and also requires the QT construction.

Our contributions are described as follows. First, we introduce the KLXX objective, whose new component mixes QT and pushforward samples. Pairwise differences remove the unknown target normalizer, and the exact-target loss retains $\nu = \mu$ as a minimizer. The target-weighted variation reuses the forward KL batch, while the mixture supplies candidate-mode and proposal-support information absent from that batch.

Second, we analyze forward KL with X_μ in Fisher–Rao geometry. The regularizer adds a nonpositive term to the instantaneous KL dissipation. For a fixed biased target surrogate and regularization strength below the stated threshold, it also improves the stationary-error bound by a nonnegative discrepancy term. These results do not cover the full KLXX loss, the changing empirical surrogate used in training, or discrete neural-network optimization.

Third, we build an adaptive-staging Boltzmann generator that uses neither external target data nor a replay buffer. AIS and QT generate the required samples on the fly, and the MCMC transitions use only derivatives of U without differentiating the flow density. Each stage compares its trained proposal with the identity on the complete validation set, retains the proposal with higher ESS, and shrinks rejected stage increments. For molecular potentials, an additional continuation sharpens a regularized interaction core toward the reported endpoint.

Fourth, we demonstrate that QT gives KLXX a mode-discovery signal unavailable to AIS-only training. QT can reach target basins outside the regions covered by the source and annealed samples. In the Himmelblau and Sparse tests, the AIS-only losses miss remote wells, while both QT-based losses recover every mode; mixing QT and pushforward samples also reduces visible intermodal leakage (Figures 3 and 4). We therefore report ESS together with coverage or a problem-specific mode diagnostic, since uniform weights on reached modes do not imply complete coverage. The clock comparison uses shared stage schedules, whereas the molecular stage schedules are method-specific and single-seed; the results support this mechanism without establishing universal or fixed-compute dominance over forward KL.

Organization. Section 2 analyzes the target-weighted regularizer, and Appendix A gives the proofs. Section 3 turns the loss into a flow-training procedure and defines its samplers and diagnostics. Section 4 uses that procedure at each stage of an adaptive-staging Boltzmann generator. Sections 5 and 6 then test the method on model and molecular targets; Section 7 concludes.

2 Fisher–Rao gradient flow analysis

We first analyze forward KL with target-weighted log-ratio variation in Fisher–Rao geometry; Appendix A gives the full derivations and proofs. This continuous-time density-space idealization exposes the mechanism of the loss but does not imply a convergence rate for the discrete optimizer: we view ν as a probability density evolving along the gradient flow and ask how fast $\text{KL}(\mu \parallel \nu_t)$ decays to zero.

For a tractable analysis we keep only the forward KL and the shape term X_μ , the log-ratio variation at weight $\omega = \mu$. The contributions weighted by $\hat{\mu}$ and $\bar{\nu}$ are left out: these fixed training measures do not enter the density-space idealization. The loss is then

$$\mathcal{L}[\nu] = \text{KL}(\mu \parallel \nu) + \lambda X_\mu(\mu \parallel \nu), \quad (9)$$

with $\lambda > 0$ fixed. Write $w = \mu/\nu$ for the importance weight and $z = \log w$ for the log-ratio. Because X_μ contains an absolute value, we use the subgradient selection $\text{sgn}(0) = 0$. The statements below apply to positive solutions of the resulting Fisher–Rao subgradient flow; for $\lambda \leq \frac{1}{2}$, its explicit birth term is nonnegative and preserves positivity.

The choice of Riemannian metric matters. In the Wasserstein-2 geometry the forward KL is not displacement convex, so the Otto–Villani/Bakry–Émery argument used for the reverse KL does not apply. We

therefore work in the Fisher–Rao (FR) geometry. Along a positive FR solution, the forward KL has an exponential bound, and X_μ adds a nonlocal birth–death term: the update at one point depends on how its importance weight ranks among weights under μ . Writing the ν -mean as $\langle \cdot \rangle_\nu := \int (\cdot) \nu dx$, the FR gradient of a smooth functional $\mathcal{F}$ and the selected flow for $\mathcal{L}$ are

$$\text{grad}^{\text{FR}} \mathcal{F}(\nu) = \nu \left(\frac{\delta \mathcal{F}}{\delta \nu} - \left\langle \frac{\delta \mathcal{F}}{\delta \nu} \right\rangle_\nu \right), \quad \partial_t \nu_t = -\text{grad}^{\text{FR}} \mathcal{L}(\nu_t), \quad (10)$$

and the mean-subtraction makes $\int \text{grad}^{\text{FR}} \mathcal{F} = 0$, so the flow conserves probability mass.

Differentiating $\text{KL}(\mu \parallel \nu_t)$ along this flow, and reducing the X_μ term by a rank-correlation identity, gives the dissipation

$$\frac{d}{dt} \text{KL}(\mu \parallel \nu_t) = -\chi^2(\mu \parallel \nu_t) - \lambda \mathbb{E}_{x, y \sim \mu} \left[\left| \frac{\mu(x)}{\nu_t(x)} - \frac{\mu(y)}{\nu_t(y)} \right| \right], \quad (11)$$

where $\chi^2(\mu \parallel \nu) := \int (\mu/\nu - 1)^2 \nu dx$ is the chi-squared divergence. The regularizer contributes a nonpositive term to the instantaneous KL dissipation. Since $\chi^2(\mu \parallel \nu_t) \geq \text{KL}(\mu \parallel \nu_t)$, Grönwall’s lemma yields

$$\text{KL}(\mu \parallel \nu_t) \leq e^{-t} \text{KL}(\mu \parallel \nu_0), \quad t \geq 0, \quad (12)$$

without assuming log-concavity, a spectral gap, or a growth condition on μ . This density-space statement does not imply that forward KL or KLXX trains faster with a discrete optimizer. It also excludes the cost of producing target samples. The added term in (11) describes how X_μ changes the density flow, not a universal transient speedup.

The convergence result assumes that the outer expectation is taken under the exact target. To isolate surrogate bias, fix a smooth density $\tilde{\mu}$ and replace μ by $\tilde{\mu}$ only in the outer integrals while retaining the exact pointwise log-ratio. This defines

$$\tilde{\mathcal{L}}[\nu] = \mathbb{E}_{x \sim \tilde{\mu}} \left[\log \frac{\mu(x)}{\nu(x)} \right] + \lambda \mathbb{E}_{x, y \sim \tilde{\mu}} \left[\left| \log \frac{\mu(x)}{\nu(x)} - \log \frac{\mu(y)}{\nu(y)} \right| \right]. \quad (13)$$

We assume that the densities are positive on their common domain and that the displayed KL, chi-squared, and pair integrals are finite. At $\lambda = 0$, the stationary density is $\nu^* = \tilde{\mu}$, so the residual error is $\text{KL}(\mu \parallel \tilde{\mu})$. The next result shows when X_μ improves this stationary-error bound. It is an idealization of one detached surrogate, not a theorem for the changing empirical batches used by the trainer.

Theorem 1 *Assume μ is absolutely continuous with respect to $\tilde{\mu}$ and $\lambda \in [0, \frac{1}{2})$. Then every stationary density ν^* of the selected subgradient flow for the fixed-surrogate loss $\tilde{\mathcal{L}}$ in (13) satisfies*

$$\text{KL}(\mu \parallel \nu^*) \leq \text{KL}(\mu \parallel \tilde{\mu}) - \lambda \iint_{\mathbb{R}^d \times \mathbb{R}^d} \tilde{\mu}(x) \tilde{\mu}(y) \left| \frac{\mu(x)}{\nu^*(x)} - \frac{\mu(y)}{\nu^*(y)} \right| dx dy. \quad (14)$$

The subtracted term is nonnegative and is positive unless the displayed importance ratio is constant on the support of $\tilde{\mu}$. Within the theorem’s assumptions, the rank feedback raises ν^* in comparatively underweighted regions and lowers it elsewhere. The experiments use $\lambda = 1$, outside the sufficient range $\lambda < \frac{1}{2}$, so the theorem does not certify their finite-training behavior. Appendix A gives the variation, rank-correlation identity, and proofs.

3 Flow training with log-ratio variation

This section turns the abstract loss of the introduction into a training algorithm for the flow map G . Let U_0 and U be the source and target potentials on $\mathbb{R}^d$, so that $\mu_0 \propto \exp(-U_0)$ is the samplable reference and $\mu \propto \exp(-U)$ is the target. We parameterize $G : \mathbb{R}^d \rightarrow \mathbb{R}^d$ so that $G_\# \mu \approx \mu_0$ and write $\nu := (G^{-1})_\# \mu_0$ for the induced approximation to μ . The loss itself evaluates only the native target-to-source map G and its log-Jacobian; it does not explicitly involve G^{-1} . The inverse map is nevertheless used to generate Monte Carlo proposals for AIS and pushforward samples for the $\bar{\nu}$ term. These are sampling-only evaluations: the sampled coordinates are detached from the inverse-map path, while their subsequent log-ratio evaluations remain differentiable through G . We first assemble the training loss and identify its sampling ingredients (Section 3.1), then supply them in turn: an AIS approximation of μ (Section 3.2), the quench and temper construction of the wide-coverage measure $\hat{\mu}$ (Section 3.3), and the ESS and coverage diagnostics used to evaluate the trained flow (Section 3.4).

3.1 Training loss with log-ratio variation

By change of variables, the pushforward density is $\nu(y) = \mu_0(G(y))|\det J_G(y)|$, so every term is a functional of the scalar field $z(y)$ in (6). The forward KL is its mean under the target, $\text{KL}(\mu\|\nu) = \mathbb{E}_{y\sim\mu}[z(y)]$. For the empirical distributions used in training, we extend (4) to any finite nonnegative measure ω by integrating against $\omega(dy)\omega(dy')$. The same formulas apply on a torus or a Euclidean-torus product with the corresponding standard reference volume. Set $s := \alpha + \beta$. For $s > 0$, let $\xi := (\alpha\hat{\mu} + \beta\bar{\nu})/s$ be the normalized mixture, with any zero-weight component omitted. Since $\alpha\hat{\mu} + \beta\bar{\nu} = s\xi$, the double integral in (4) produces the factor s^2 below. The map-level loss is

$$\mathcal{L}[G] = \underbrace{\mathbb{E}_{y\sim\mu}[z(y)]}_{\text{KL}(\mu\|\nu)} + \lambda \underbrace{\mathbb{E}_{y,y'\sim\mu}[|z(y) - z(y')|]}_{X_\mu(\mu\|\nu)} + s^2 \underbrace{\mathbb{E}_{y,y'\sim\xi}[|z(y) - z(y')|]}_{X_{\alpha\hat{\mu}+\beta\bar{\nu}}(\mu\|\nu)}. \quad (15)$$

When $s = 0$, the last term is omitted. Equation (15) is the map form of (5) after removing the introductory restriction $s = 1$; the factor s^2 follows from the double integral in (4) and is required when $\alpha + \beta \neq 1$. The default choice $\lambda = 1$ and $\alpha = \beta = \frac{1}{2}$ defines KLXX. The pushforward distribution $\nu = (G^{-1})_{\#}\mu_0$ depends on the trainable parameters of G . The barred measure $\bar{\nu}$ is the same distribution frozen at the current parameters: samples are drawn through G^{-1} and detached from that sampling path, but z at the sampled coordinates is still differentiated through the native map G .

The three sampling measures entering (15) address complementary failure modes of the bare forward KL:

- μ : *target accuracy*. The term X_μ penalizes pairwise log-ratio variation within the target batch. Theorem 1 gives its density-space bound for $\lambda < \frac{1}{2}$; the finite-training effect at $\lambda = 1$ is tested empirically.
- $\hat{\mu}$: *candidate mode coverage*. QT provides comparisons on basins represented in its wide-coverage set, including basins absent from the annealed batch. This supplies a training signal for those represented basins but does not guarantee that QT finds every mode.
- $\bar{\nu}$: *proposal support*. Pushforward samples evaluate the log-ratio where the current proposal places mass. The term is intended to reduce variation in target-poor regions; the experiments test whether this is accompanied by less visible leakage.

Adding these nonnegative terms to the exact-target forward KL preserves $\nu = \mu$ as a minimizer. The terms change the sampled training signal without changing that reference solution.

For $B \geq 2$, the forward KL and X_μ share one target batch: evaluate $\mathbf{z} = (z(y_1), \dots, z(y_B))$, take $\text{mean}(\mathbf{z})$ for the forward KL, and $\text{mean}(|\mathbf{z} - \mathbf{z}[\sigma]|)$ for a uniformly random derangement σ . This is an unbiased randomized estimate of the distinct-pair average and adds only a pairing and pairwise difference to the same differentiation graph. When $s > 0$, the last term evaluates one additional batch drawn from ξ , applies the same randomized pairwise estimator, and multiplies it by s^2 . At the KLXX setting, $s = 1$ and the batch is split equally between $\hat{\mu}$ and $\bar{\nu}$.

The loss draws on three measures, of which at most two must be built. The $\bar{\nu}$ batch is the raw pushforward $\bar{y}_i = G^{-1}(x_i)$ of source samples $x_i \sim \mu_0$. The inverse-map sampling path is detached, while $z(\bar{y}_i)$ is then evaluated through the native map G and differentiated normally. An approximation of the target μ for the forward KL and X_μ terms is supplied by the annealing procedure of Section 3.2. When $\alpha > 0$, a second construction supplies samples of the wide-coverage measure $\hat{\mu}$ by quench and temper (Section 3.3). Throughout loss differentiation, the sampled coordinates and all resampling, MALA, mixture, and derangement choices are treated as fixed; gradients pass only through evaluations of z at the resulting coordinates. With equal target and mixture batch sizes, $\text{KL} + X_\mu + X_{(\hat{\mu} + \bar{\nu})/2}$ evaluates about twice as many samples per optimizer iteration as forward KL. This does not imply twice the total training time: constructing $\hat{\mu}$ also applies quench and temper to the complete validation set and adds a problem- and stage-dependent cost. Inference uses the same flow architecture for both methods.

Two recurring settings are the X_μ -only setting, with $\lambda = 1$ and $\alpha = \beta = 0$, and KLXX, with $\lambda = 1$ and $\alpha = \beta = \frac{1}{2}$. The first adds no sampling measure beyond the target batch, while the second splits its mixture batch equally between $\hat{\mu}$ and $\bar{\nu}$. We report the full KLXX loss as $\text{KL} + X_\mu + X_{(\hat{\mu} + \bar{\nu})/2}$ when it must be distinguished from the other ablations.

Remark 1 (Why forward KL is still necessary) *The X_μ -only setting above means $KL + X_\mu$; it does not remove forward KL. Log-ratio variation alone is not a reliable training objective: empirical pairwise terms constrain relative log-ratios only on sampled regions and provide no global density-shaping constraint outside them. Appendix D.2 of the LDR study (Schopmans et al., 2026) confirms the analogous failure: without a data-based divergence, training with only LDR-1 or only LDR-2 produces qualitatively incorrect distributions. We therefore always combine X-regularization with forward KL.*

3.2 Approximating μ by annealed importance sampling

Evaluating the μ -expectations in (15) requires samples from μ , which are not directly available. We form a surrogate by drawing Monte Carlo proposals $y_1, \dots, y_B$ from the pushforward $\nu = (G^{-1})_\# \mu_0$ and reweighting them. In the theory, the normalized importance ratio is $w = \mu/\nu$. Computation uses the proportional weight

$$\tilde{w}(y) := \exp \left(U_0(G(y)) - U(y) - \log |\det J_G(y)| \right) \propto \frac{\mu(y)}{\nu(y)}, \quad y \in \mathbb{R}^d, \quad (16)$$

Below, w_i denotes these proportional weights because resampling, normalized ESS, and log-ratio differences are scale invariant. A resampling step and a short Metropolis-adjusted Langevin algorithm (MALA) simulation then diffuses the duplicates. In high dimensions, however, μ and ν can differ greatly, so the weights can collapse onto a few outliers. We temper the reweighting through $M + 1$ intermediate distributions

$$\pi_m(y) \propto \nu(y)^{1-\frac{m}{M}} \mu(y)^{\frac{m}{M}} \propto ((G^{-1})_\# \mu_0)(y) \tilde{w}(y)^{\frac{m}{M}}, \quad m = 0, 1, \dots, M, \quad (17)$$

so that $\pi_0 = \nu$ and $\pi_M = \mu$. Each step reweights by only $\tilde{w}(y)^{1/M}$, reducing the degeneracy of direct importance sampling.

Algorithm 1: Annealed importance sampling with pushforward proposals toward μ

Input: Monte Carlo proposals $y_1, \dots, y_B$ from $(G^{-1})_\# \mu_0$, AIS ladder length M , intermediate distributions π_m (17)

Output: biased target-surrogate batch $y_1, \dots, y_B$ approximating μ

```

1 for  $m = 1, \dots, M$  do
2   Compute the incremental weights  $\tilde{w}(y_i)^{1/M} = \exp(z(y_i)/M)$  on  $y_1, \dots, y_B$ ;
   // resample
3   Resample from the weighted set,  $y'_1, \dots, y'_B \leftarrow \text{resample}(y_{1:B}, (\tilde{w}(y_i)^{1/M})_{i=1}^B)$ ;
   // rejuvenate under the target
4   Diffuse the resampled duplicates by a short MALA chain on  $\mu$ ,  $y_{1:B} \leftarrow \text{MALA}(y'_{1:B}, U)$ ;
5 end
```

An exact SMC sampler on the path (17) would rejuvenate at step m with a Markov kernel invariant for π_m . A MALA kernel for π_m would require $\nabla_y \log \nu(y)$ and hence positional derivatives of the flow density, including $\nabla_y \log |\det J_G(y)|$. Algorithm 1 instead applies MALA directly under μ , using only $-\nabla U$ in its proposal. MALA preserves μ , but not the intermediate π_m , so the output is a biased target surrogate rather than exact samples from μ . The construction trades intermediate invariance for a sampler that moves resampled particles under the final target. Section 2 studies a fixed smooth surrogate as a density-space idealization; it does not model the changing empirical output of Algorithm 1 exactly. Evaluation therefore uses fresh importance weights and coverage diagnostics.

Remark 2 (Choosing the SMC/AIS ladder length M) *For a target close to the source, $M = 1$ can suffice. When the change is distributed across many coordinates, the log-weight increment accumulates their contributions, so M should grow with the effective dimension of the deformation. The choice $M = d$ is a conservative default when all d coordinates change; a shorter ladder is appropriate when only a few directions differ.*

Remark 3 (Comparison with FAB) *FAB (Midgley et al., 2022) also couples AIS to flow training, but places the flow inside the chains. Its MCMC transitions target the surrogate μ^2/ν and require the log-density gradient*

$$\nabla_y \log \frac{\mu^2(y)}{\nu(y)} = -2\nabla U(y) - \nabla_y \log \nu(y)$$

which contains the positional gradient of the pushforward density at every MCMC step. Training uses a prioritized replay buffer of weighted AIS samples with a staleness correction. Reusing the buffer reduces repeated AIS generation but adds storage, prioritized sampling, and correction work. Here every MALA chain targets μ and uses only $-\nabla U$; the flow enters only through the initial G^{-1} proposal and evaluations of z . Each gradient step starts from a new source batch, so no replay buffer is kept. Relative to forward KL, KLXX instead adds the QT construction and pairwise log-ratio variation evaluations: X_μ reuses the target batch, while the mixture term evaluates an additional batch. The fixed N -sample QT set of Section 3.3 supplies separate coverage information. FAB instead relies on its AIS budget for exploration. The 2-divergence gives FAB a stronger tail-coverage incentive, and with sufficient AIS work it may need no separate coverage construction.

3.3 Constructing $\hat{\mu}$ by quench and temper

When $\alpha > 0$, the last term needs samples from the wide-coverage measure $\hat{\mu}$. Algorithms 2 and 4 use the complete validation set $\mathcal{Y} = \{x_1, \dots, x_N\}$ for each stage, and QT acts on all N samples once per stage attempt. The resulting set is held fixed as a reservoir during optimization. At each training step a mini-batch of B samples is drawn uniformly from that reservoir and refreshed by a short MALA chain before entering the X_ω computation. We use $\hat{\mu}$ for the law of this draw-and-refresh operation; its input is the uniform empirical measure on the fixed QT reservoir. Standalone one-flow ablations may instead report a separately sized QT set; the staged algorithm uses only the validation set size N and batch size B . Quench and temper (QT) is the map

$$\hat{\mathcal{Y}} = \text{QT}(\mathcal{Y}, U)$$

from the validation set and target potential U to a fixed wide-coverage reservoir $\hat{\mathcal{Y}} = \{\hat{x}_1, \dots, \hat{x}_N\}$. Its uniform empirical distribution is the input to the per-step refresh defining $\hat{\mu}$. QT consists of melt, quench, and temper.

Algorithm 2: Quench and temper (QT): wide-coverage construction

Input: validation set $\mathcal{Y} = \{x_1, \dots, x_N\}$, target potential U , melt scale m_e
Output: wide-coverage set $\hat{\mathcal{Y}} = \{\hat{x}_1, \dots, \hat{x}_N\}$ forming $\hat{\mu}$, written $\hat{\mathcal{Y}} = \text{QT}(\mathcal{Y}, U)$
`// melt`
1 Scatter the samples $x_i \leftarrow x_i + m_e \zeta_i$ with $\zeta_i \sim \mathcal{N}(0, I_d)$ i.i.d. and melt scale $m_e > 0$;
`// quench`
2 Drive each sample toward a mode center of μ by L-BFGS on U , updating $x_1, \dots, x_N$;
`// temper`
3 Spread the samples around each mode, $\hat{x}_{1:N} \leftarrow \text{MALA}(x_{1:N}, U)$;

The quench places samples near local minima of U , and the Gaussian scatter broadens the set of basins that can be reached. On a torus or Euclidean-torus product, periodic coordinates are interpreted modulo their period after the melt and after each MALA proposal. If the QT set contains a basin absent from the annealed target batch, $X_{\hat{\mu}}$ supplies pairwise log-ratio evaluations on that basin. QT does not certify complete mode discovery.

A large melt scale m_e can bias the QT reservoir toward basins far from the input set. This is intentional: $\hat{\mu}$ is a coverage measure, not an estimate of target mode probabilities. Adding $X_{\hat{\mu}}$ to the exact-target forward KL preserves $\nu = \mu$ as a minimizer, while empirical frequencies in the finite reservoir affect the optimization path.

3.4 Importance sampling at evaluation: ESS and coverage

At evaluation time the trained flow G supplies the pushforward distribution $\nu = (G^{-1})_{\#}\mu_0$ and its tractable density $\nu(y) = \mu_0(G(y))|\det J_G(y)|$. The proportional importance weight $\tilde{w}(y)$ in (16) can therefore be evaluated on any pushforward sample without fitting another density model. Draw $\mathbf{y} = (y_1, \dots, y_N)$ with $y_i = G^{-1}(x_i)$ and $x_i \sim \mu_0$ independently, and write $\mathbf{w} = (w_1, \dots, w_N)$ for any common scaling of the corresponding positive weights. These weights diagnose agreement on reached support and form self-normalized importance estimates or a reweighted or resampled target approximation, followed when desired by a short MALA refresh under μ .

We assess this pipeline through two metrics: the effective sample size, which measures how uniform the importance weights are on the support reached by the flow, and the coverage, which checks whether the regions represented in the QT reference set have been reached.

The first metric is the effective sample size, normalized to $[0, 1]$ by dividing the classical ESS by the sample count N . For the weight vector $\mathbf{w}$, define

$$\text{ESS}(\mathbf{w}) = \frac{(\sum_{i=1}^N w_i)^2}{N \sum_{i=1}^N w_i^2} \in [0, 1]. \quad (18)$$

The upper bound follows from Cauchy–Schwarz and is attained when all weights are equal. The factor N distinguishes this normalized ESS from the classical count $(\sum_i w_i)^2 / \sum_i w_i^2 \in [1, N]$ (Elvira et al., 2022). Multiplying all weights by a common positive constant leaves both definitions unchanged. A high ESS indicates uniform observed weights only on support reached by ν ; a target mode absent from the sample cannot lower it.

To detect such missing modes, we pair the ESS with the coverage metric of Naeem et al. (2020). Let $\hat{\mathbf{y}} = (\hat{y}_1, \dots, \hat{y}_N) \in \mathbb{R}^{N \times d}$ be the fixed QT reference reservoir produced by Algorithm 2. For each $1 \leq i \leq N$, the κ -nearest-neighbor radius of $\hat{y}_i$ (with $\kappa = 5$) is

$$\text{NND}_{\kappa}(\hat{y}_i; \hat{\mathbf{y}}) = \inf \left\{ r \geq 0 : \#\{1 \leq j \leq N : 0 < |\hat{y}_i - \hat{y}_j| < r\} \geq \kappa \right\}, \quad (19)$$

i.e., the distance from $\hat{y}_i$ to its κ -th nearest neighbor in $\hat{\mathbf{y}}$. The coverage of $\hat{\mathbf{y}}$ by the pushforward samples $\mathbf{y}$ is then

$$\text{Coverage}_{\kappa}(\mathbf{y}; \hat{\mathbf{y}}) = \frac{1}{N} \sum_{i=1}^N \mathbf{1} \left\{ \exists 1 \leq j \leq N : |y_j - \hat{y}_i| \leq \text{NND}_{\kappa}(\hat{y}_i; \hat{\mathbf{y}}) \right\}, \quad (20)$$

i.e., the fraction of points in $\hat{\mathbf{y}}$ whose κ -nearest-neighbor ball contains at least one pushforward sample. A coverage close to 1 indicates that the finite QT reference set is well covered by the flow, while a coverage well below 1 exposes missing regions that the ESS alone cannot see. Because the reference reservoir is itself produced by QT, the diagnostic checks only modes found by that reference construction; it cannot reveal a mode absent from the reference set.

4 Adaptive-staging Boltzmann generator

A single flow can be inadequate when the source-to-target change is too large, so we divide the interpolation into stages. Each stage changes the potential by a controlled amount and trains one flow with the loss in (15), using user-chosen weights $\lambda, \alpha, \beta \geq 0$. The KLXX Boltzmann generator is the special case $(\lambda, \alpha, \beta) = (1, \frac{1}{2}, \frac{1}{2})$; the construction below also permits other choices. Sampling diagnostics determine the stage schedule.

4.1 Construction of the adaptive stage schedule

To bridge the simple source $\mu_0 \propto \exp(-U_0)$ to the target $\mu \propto \exp(-U)$, we seek an increasing sequence

$$0 = t_0 < t_1 < \dots < t_K = 1$$

defining the interpolating potential and measure

$$U_t(x) = (1-t)U_0(x) + tU(x), \quad \mu_t(x) \propto \exp(-U_t(x)), \quad t \in [0, 1].$$

Here t is a dimensionless interpolation parameter, not a thermodynamic temperature. The discrete values t_k are stage points, and stage k is the transition from t_{k-1} to t_k . At the accepted stage points, write $\mu_k := \mu_{t_k}$, so μ_0 is the source and μ_K is the target. Stage k trains a target-to-source map G_k , whose inverse induces the proposal $\nu_k = (G_k^{-1})_{\#}\mu_{k-1}$ for μ_k . The number of stages K is selected during the run. Its central choice is the stage increment $\Delta t_k = t_k - t_{k-1}$: a large increment can be hard to fit, whereas a small increment adds stages.

Training G_k requires the target surrogate for μ_k defined in Section 3.2. Before training, an optional classical SMC calculation can screen the proposed stage increment Δt_k . Its smallest incremental ESS measures overlap along that candidate stage for the current particle set. This is a difficulty diagnostic, not a mode-coverage guarantee. Its floor lies in $[0, 1)$, and setting it to zero disables the screen.

Algorithm 3: Adaptive selection of the stage increment Δt_k

Input: potentials U_0, U , previous stage point t_{k-1} , validation set $\mathcal{Y}_{k-1} = \{y_1, \dots, y_N\}$
approximating μ_{k-1} , SMC ladder length M , SMC ESS floor $\tau_{\text{SMC}} = 0$, shrink factor
 $\gamma = 0.7$, initial guess t_{init}

Output: next stage point t_k , giving the stage increment $\Delta t_k = t_k - t_{k-1}$

```

1  $U_{t_{k-1}} \leftarrow (1 - t_{k-1})U_0 + t_{k-1}U;$ 
  // start from the supplied initial guess
2  $t_k \leftarrow t_{\text{init}};$ 
3 if  $\tau_{\text{SMC}} = 0$  then
4   | return  $t_k;$ 
5 end
6 while true do
  // run a classical SMC ladder across the candidate stage
7   Run an  $M$ -step classical SMC ladder on a working copy of all  $N$  particles in  $\mathcal{Y}_{k-1}$ , using the
    linear potential bridge from  $U_{t_{k-1}}$  to  $U_{t_k}$  and a MALA kernel invariant for each intermediate
    distribution ( $M$  as in Remark 2);
  // evaluate the ladder by its worst per-step ESS
8   Record  $\text{ESS}_{\min}$ , the smallest per-step ESS over the  $M$  ladder steps;
9   if  $\text{ESS}_{\min} \geq \tau_{\text{SMC}}$  then
10    | return  $t_k;$ 
11  else
12    | // shrink the stage increment toward  $t_{k-1}$ 
13    |  $t_k \leftarrow t_{k-1} + \gamma(t_k - t_{k-1});$ 
14  end
15 end

```

When enabled, these SMC ladders involve only the potentials U_0 and U , never the flow map. They use no QT samples, so their per-step ESS can overlook a mode the SMC samples never reach. We therefore use this ESS only to screen a candidate stage; it is a heuristic measure of stage difficulty, not a coverage certificate. The selected-proposal validation in step (v) supplies the acceptance gate, while QT supplies the separate wide-coverage set when $\alpha > 0$. By default $\tau_{\text{SMC}} = 0$, so Algorithm 3 returns t_{init} without running SMC. For a positive floor, it returns the first candidate stage along the shrink sequence whose estimated ESS clears that floor. The shrink factor is $\gamma = 0.7$. As t_k approaches t_{k-1} , the stage becomes trivial and its ESS approaches 1, so a floor below 1 is attainable under repeated shrinking in exact arithmetic.

4.2 The full training procedure

Algorithm 4 combines the stage-selection rule with the training procedure. It carries a fixed-size validation set $\mathcal{Y}_k \approx \mu_k$ across stages. Each candidate stage builds its QT set when $\alpha > 0$, trains a flow with the selected loss weights, and compares that flow with the identity proposal. If neither proposal reaches the validation ESS floor $\tau_v \in [0, 1)$, the algorithm shrinks the stage increment and retries. The displayed loop ends when $t_K = 1$. Practical implementations use an endpoint tolerance and finite caps on retries and stages; reaching a cap returns an incomplete stage schedule.

- (i) *Validation set.* Stage k uses the *validation set size* N and the *batch size* B . The validation set $\mathcal{Y}_{k-1} = \{y_1, \dots, y_N\} \approx \mu_{k-1}$ is carried over from the previous stage. A working copy of all N samples supplies the SMC calculation in step (ii), and, when $\alpha > 0$, all N samples initialize QT in step (iii). Neither operation first draws a separately sized subset. Training batches are drawn at run time, B samples per gradient step, from $\mathcal{Y}_{k-1}$; no separate training set is built.
- (ii) *Adaptive stage selection.* Algorithm 3 selects the next stage point t_k . The first stage starts from $t_{\text{init}} = t_{\text{safe}} = 0.2$. Later stages extrapolate the last accepted stage increment, $t_{\text{init}} = \min(1, t_{k-1} + \Gamma(t_{k-1} - t_{k-2}))$, with enlarge factor $\Gamma = 1.5$. The default $\tau_{\text{SMC}} = 0$ accepts this initial guess without the optional SMC screen. If a positive floor is supplied, the algorithm shrinks the stage increment until the classical SMC ESS on all N particles clears that floor. The selected stage target is $\mu_k \propto \exp(-U_{t_k})$, with $U_{t_k} = (1 - t_k)U_0 + t_k U$.
- (iii) *Wide-coverage set.* When $\alpha > 0$, quench and temper (Algorithm 2) acts on the complete validation set $\mathcal{Y}_{k-1}$ under U_{t_k} . The resulting N samples $\hat{\mathcal{Y}}_k = \text{QT}(\mathcal{Y}_{k-1}, U_{t_k})$ form the fixed stage reservoir. The measure $\hat{\mu}_k$ is the law obtained by drawing uniformly from this reservoir and applying the short MALA refresh used at each gradient step; it supplies the wide-coverage component of the last term in (21). When $\alpha = 0$, that component vanishes and the QT construction is omitted.
- (iv) *Training loss.* The flow G_k is trained by the general loss (15) specialized to the stage $\mu_{k-1} \rightarrow \mu_k$. With the pushforward distribution $\nu_k = (G_k^{-1})_{\#} \mu_{k-1}$ and its frozen sampling measure $\bar{\nu}_k$, the stage log-ratio is

$$z_k(y) = U_{t_{k-1}}(G_k(y)) - U_{t_k}(y) - \log |\det J_{G_k}(y)|.$$

Let $s = \alpha + \beta$. For $s > 0$, define the stage specialization of the normalized mixture ξ , with any zero-weight component omitted,

$$\xi_k := \frac{\alpha \hat{\mu}_k + \beta \bar{\nu}_k}{s}.$$

The stage loss then reads

$$\mathcal{L}[G_k] = \underbrace{\mathbb{E}_{y \sim \mu_k} [z_k(y)]}_{\text{KL}(\mu_k \parallel \nu_k)} + \lambda \underbrace{\mathbb{E}_{y, y' \sim \mu_k} [|z_k(y) - z_k(y')|]}_{X_{\mu_k}(\mu_k \parallel \nu_k)} + s^2 \underbrace{\mathbb{E}_{y, y' \sim \xi_k} [|z_k(y) - z_k(y')|]}_{X_{\alpha \hat{\mu}_k + \beta \bar{\nu}_k}(\mu_k \parallel \nu_k)}, \quad (21)$$

When $s = 0$, the last term is omitted. The μ_k -expectations are estimated from the biased AIS surrogate: a batch $\mathcal{X}$ drawn from $\mathcal{Y}_{k-1}$ is pushed to the proposal batch $\bar{\mathcal{X}}_\nu = G_k^{-1}(\mathcal{X})$, approximating ν_k , and annealed toward μ_k by Algorithm 1 with $G = G_k$. When $s > 0$, a batch $\mathcal{X}_{\xi_k}$ from ξ_k is formed by drawing from $\hat{\mu}_k$ with probability α/s when $\alpha > 0$, and from $\bar{\nu}_k$ with probability β/s when $\beta > 0$; the randomized pairwise estimate is then multiplied by s^2 . This factor retains the general-weight definition in (4) when $\alpha + \beta \neq 1$. The choices $(\lambda, \alpha, \beta) = (0, 0, 0)$ and $(1, 0, 0)$ give forward KL and the X_{μ_k} -only loss, respectively; $(1, \frac{1}{2}, \frac{1}{2})$ gives the default KLXX loss.

- (v) *Validation set update.* Once the candidate G_k is trained, each $y \in \mathcal{Y}_{k-1}$ is pushed through its inverse to $\tilde{y} = G_k^{-1}(y)$. This gives the trained proposal set $\tilde{\mathcal{Y}}_k^{\text{tr}} = G_k^{-1}(\mathcal{Y}_{k-1})$ from $\nu_k = (G_k^{-1})_{\#} \mu_{k-1}$, together with proportional importance weights represented by

$$w_k^{\text{tr}}(\tilde{y}) := \exp(z_k(\tilde{y})), \quad \tilde{y} \in \tilde{\mathcal{Y}}_k^{\text{tr}}.$$

The identity proposal keeps $\mathcal{Y}_{k-1}$ unchanged and has weights $w_k^{\text{id}}(y) = \exp[U_{t_{k-1}}(y) - U_{t_k}(y)]$. We select the pair with the larger ESS over the complete validation set and denote it by $(\tilde{\mathcal{Y}}_k, w_k)$, taking

the stage map to be G_k or the identity accordingly. If $\text{ESS}(w_k) < \tau_v$, with default $\tau_v = 0.6$, we discard the candidate, shrink to $t_k \leftarrow t_{k-1} + \gamma(t_k - t_{k-1})$, and repeat steps (iii)–(v) at the reduced stage increment. Otherwise we advance the validation set to $\mathcal{Y}_k = \text{MALA}(\text{resample}(\hat{\mathcal{Y}}_k, w_k; N), U_{t_k}) \approx \mu_k$.

The accepted stage schedule defines a Boltzmann generator for μ . Starting with samples from μ_0 , stage k applies the selected inverse map or identity, computes w_k , resamples, and rejuvenates under U_{t_k} . The resulting validation set supplies stage $k + 1$.

The accepted-stage weights also define the propagation factor used below:

$$F := \prod_{k=1}^K \text{ESS}(w_k)^{-1}. \quad (22)$$

Because normalized ESS lies in $(0, 1]$, F is at least one, with the ideal value one. It summarizes stagewise weight degeneracy by multiplying reciprocal incremental ESS values. It is neither the ESS of the complete composed sampler nor an estimator of endpoint Monte Carlo error. If a stage contains another accepted reweighting operation, such as molecular sharpening in Section 6, its reciprocal ESS enters the product.

Algorithm 4: Adaptive-staging Boltzmann generator training

Input: potentials U_0, U ; sizes $N, B \geq 2$; loss weights $\lambda, \alpha, \beta \geq 0$; SMC/AIS ladder length M ; ESS floors $\tau_{\text{SMC}} = 0, \tau_v = 0.6$; factors $\gamma = 0.7, \Gamma = 1.5$; safe start $t_{\text{safe}} = 0.2$

Output: completed stage schedule $0 = t_0 < \dots < t_K = 1$ and selected stage maps $\{G_k\}_{k=1}^K$

// validation set, $\mathcal{Y}_k \approx \mu_k$

```
1  $\mathcal{Y}_0 \leftarrow N$  samples from  $\mu_0$ ,  $t_0 \leftarrow 0$ ,  $k \leftarrow 0$ ;  
2 while  $t_k < 1$  do  
3    $k \leftarrow k + 1$ ;  
4   Set  $U_{t_{k-1}} \leftarrow (1 - t_{k-1})U_0 + t_{k-1}U$ ;  
   // (i) validation set  
5   Use  $\mathcal{Y}_{k-1}$  for stage selection, validation, and QT when  $\alpha > 0$ ;  
   // (ii) adaptive stage selection  
6   if  $k = 1$  then  
7     Set  $t_{\text{init}} = t_{\text{safe}}$ ;  
8   else  
9     Set  $t_{\text{init}} = \min(1, t_{k-1} + \Gamma(t_{k-1} - t_{k-2}))$ ;  
10  end  
11  Choose  $t_k$  by Algorithm 3 on  $\mathcal{Y}_{k-1}$  from  $t_{\text{init}}$ ;  
12  Set  $U_{t_k} \leftarrow (1 - t_k)U_0 + t_kU$ ;  
13  repeat  
14    Initialize  $G_k$  to the identity map;  
    // (iii) optional wide-coverage set  
15    If  $\alpha > 0$ , set  $\hat{\mathcal{Y}}_k \leftarrow \text{QT}(\mathcal{Y}_{k-1}, U_{t_k})$  by Algorithm 2;  
    // (iv) general log-ratio-variation loss  
16    for each gradient step do  
17      Draw  $B$  samples  $\mathcal{X}$  from  $\mathcal{Y}_{k-1}$ ;  
18      Set the pushforward batch  $\tilde{\mathcal{X}}_\nu \leftarrow G_k^{-1}(\mathcal{X})$  and detach its sampling path;  
19      Anneal these Monte Carlo proposals toward  $\mu_k$  to form the biased target batch  $\mathcal{X}_\mu$  (Algorithm 1);  
20      If  $\alpha > 0$ , draw  $\hat{\mathcal{X}}_\mu \subset \hat{\mathcal{Y}}_k$  and refresh it by MALA under  $U_{t_k}$ ;  
21       $s \leftarrow \alpha + \beta$ ; if  $s > 0$ , mix  $\hat{\mathcal{X}}_\mu$  (when  $\alpha > 0$ ) and  $\tilde{\mathcal{X}}_\nu$  with probabilities  $\alpha/s, \beta/s$  to form  $\mathcal{X}_{\xi_k}$ ;  
22      Update  $G_k$  with the loss below using random derangements; omit its last term if  $s = 0$ :  
  
      
$$\hat{\mathcal{L}}[G_k] = \hat{\mathbb{E}}_{y \in \mathcal{X}_\mu} [z_k(y)] + \lambda \hat{\mathbb{E}}_{y, y' \in \mathcal{X}_\mu} [|z_k(y) - z_k(y')|] + s^2 \hat{\mathbb{E}}_{y, y' \in \mathcal{X}_{\xi_k}} [|z_k(y) - z_k(y')|].$$
  
23    end  
    // (v) validation set update  
24    Form the trained and identity pairs as in step (v);  
25    Keep the larger-ESS pair as  $(\tilde{\mathcal{Y}}_k, w_k)$  and set  $G_k$  to the corresponding map;  
26    if  $\text{ESS}(w_k) < \tau_v$  then  
27      // reject the stage and shrink toward  $t_{k-1}$   
       $t_k \leftarrow t_{k-1} + \gamma(t_k - t_{k-1})$ , then update  $U_{t_k} \leftarrow (1 - t_k)U_0 + t_kU$ ;  
28    else  
29       $\mathcal{Y}_k \leftarrow \text{MALA}(\text{resample}(\tilde{\mathcal{Y}}_k, w_k; N), U_{t_k})$ ;  
30    end  
31  until  $\text{ESS}(w_k) \geq \tau_v$ ;  
32 end  
33  $K \leftarrow k$ ;  
34 return  $\{t_k\}_{k=0}^K$  and  $\{G_k\}_{k=1}^K$ ;
```

5 Numerical experiments on model problems

We now test the roles assigned to the three sampling distributions. The 2D targets assess target accuracy, QT-assisted mode coverage, and visible proposal leakage. The product multi-well tests how the ESS difference changes with dimension. The tilted ϕ^4 model shows why ESS must be paired with a collective-mode diagnostic. The clock model then compares the full objective on shared annealing schedules. Accordingly, ESS is reported with coverage, minority-phase weight, or sector occupancy whenever mode omission is possible.

5.1 2D benchmark distributions

We consider four 2D targets: Two-Moon, Three-Well, Himmelblau, and Sparse. Two-Moon has two distorted crescent components and isolates finite-batch accuracy because both components lie within reach of the source. The remaining targets isolate mode collapse: Three-Well contains a low-probability basin, Himmelblau has four separated wells, and Sparse places two tight modes far into the source tails.

For every target, we initialize the same neural spline flow at the identity and compare four objectives:

$$\text{KL}, \quad \text{KL}+X_\mu, \quad \text{KL}+X_\mu+X_{\hat{\mu}}, \quad \text{KL}+X_\mu+X_{(\hat{\mu}+\bar{\nu})/2}.$$

The flow has 32 spline bins, 6 transforms, and conditioner widths (128, 128). Each training step constructs a target surrogate by one AIS update through the current flow. The rejuvenation kernel uses 100 MALA steps of size 10^{-3} for Two-Moon and 50 MALA steps of size 2×10^{-3} for the other targets. The fixed validation sets contain 5×10^4 points for Two-Moon and Himmelblau and 8×10^4 for Three-Well and Sparse. For both objectives containing the wide-coverage term $X_{\hat{\mu}}$, quench and temper acts on the complete fixed validation set, with melt scale 2 and an Armijo L-BFGS quench using initial step 0.5 and 200 iterations. The final ESS is computed from flow importance weights on that full validation set. Coverage, with $\kappa = 5$ against an independently generated quench and temper reference set, measures the fraction of reference configurations represented by the pushforward through the nearest-neighbor test of Section 3.4. The reference sets contain 1000, 4000, 500, and 4000 points for Two-Moon, Three-Well, Himmelblau, and Sparse, respectively; their QT tempering uses 1000 MALA steps at the target-specific step size above. Hence a collapsed flow can have ESS near one but cannot have full coverage. Each objective–target pair is one fixed-seed training realization, so the numerical differences below describe these matched runs rather than variation across independent trainings.




Figure 1: Two-Moon target. Columns, from left to right, are forward KL, $\text{KL}+X_\mu$, $\text{KL}+X_\mu+X_{\hat{\mu}}$, and $\text{KL}+X_\mu+X_{(\hat{\mu}+\bar{\nu})/2}$. Each panel overlays the learned pushforward on the target energy and reports its final ESS and coverage. Gray points show the source.

Two-Moon. The source is $\mathcal{N}(0, I_2)$ and both crescents lie within its reach, so all four losses attain coverage 1. With batch size 500, 200 Adam steps, and learning rate 10^{-3} , the bare forward KL reaches ESS 0.8358. Adding X_μ raises it to 0.9109, while the equal-weight mixture is best at 0.9236 (Figure 1). Because coverage is already complete, this comparison tests whether pairwise log-ratio control improves density calibration on the reached support.




Figure 2: Three-Well target. From left to right, the panels show forward KL, $\text{KL}+X_\mu$, $\text{KL}+X_\mu+X_{\hat{\mu}}$, and $\text{KL}+X_\mu+X_{(\hat{\mu}+\bar{\nu})/2}$. Each panel overlays the learned pushforward on the target energy and reports final ESS and coverage; gray points show the source. Forward KL and $\text{KL}+X_\mu$ represent the two favored wells but miss the suppressed third well. Both objectives weighted by $\hat{\mu}$ place samples in all three wells, and the equal mixture retains the higher ESS of the two mode-complete fits.

Three-Well. The source is $\mathcal{N}(0, 0.25^2 I_2)$, narrower than the well separation and initially covering no basin. Training uses batch size 1000, 500 steps, and learning rate 10^{-3} . Forward KL and $\text{KL}+X_\mu$ fill the two favored wells and report ESS 0.9884 and 0.9902, respectively, while their coverage remains 0.7595. The QT set also represents the suppressed third well. Both objectives that use it cover all three wells in the sample panels, with scalar coverages 0.9783 and 0.9768 and ESS 0.9611 and 0.9893 (Figure 2).




Figure 3: Himmelblau target. Columns show forward KL, $\text{KL}+X_\mu$, $\text{KL}+X_\mu+X_{\hat{\mu}}$, and $\text{KL}+X_\mu+X_{(\hat{\mu}+\bar{\nu})/2}$; each panel shows the pushforward over the target energy with final ESS and coverage, and gray points show the source. The first two objectives capture only two of the four wells. Both $\hat{\mu}$ -weighted objectives cover all four, while including pushforward samples in the mixture term reduces the visible intermodal mass and raises ESS relative to the $\hat{\mu}$ -only fit.

Himmelblau. The source is $\mathcal{N}(0, I_2)$, whereas the four wells lie at radius approximately 3–4. Training uses batch size 1000, 1000 steps, and learning rate 10^{-3} . Forward KL and $\text{KL}+X_\mu$ capture only two wells, giving coverage 0.5160 despite ESS 0.9316 and 0.9809. Both objectives using $\hat{\mu}$ attain coverage 1. Relative to the $\hat{\mu}$ -only result, the equal mixture shows less intermodal overcoverage and raises ESS from 0.9562 to 0.9736 (Figure 3). This paired observation is consistent with, but does not by itself prove, the proposed leakage-control role of $\bar{\nu}$.




Figure 4: Sparse target. Columns show forward KL, $\text{KL}+X_\mu$, $\text{KL}+X_\mu+X_{\hat{\mu}}$, and $\text{KL}+X_\mu+X_{(\hat{\mu}+\bar{\nu})/2}$, with the learned pushforward overlaid on the target energy; panel titles report final ESS and coverage, and gray points show the source. The first two objectives cover the central pair but miss the two remote modes. The $\hat{\mu}$ -only objective restores the remote modes with visible mass between components, whereas the equal mixture covers all four modes with less intermodal mass and higher ESS.

Sparse. The target is a four-component Gaussian mixture with component standard deviation 0.08: two modes lie near the source center and two lie in opposite corners, more than five source standard deviations away. With source standard deviation 0.6, batch size 1000, 1000 steps, and learning rate 5×10^{-3} , forward KL and $\text{KL}+X_\mu$ cover only the central pair. Their coverage is 0.5335, while their ESS is 0.9116 and 0.9266. The $\hat{\mu}$ -only objective restores all four modes but also places visible mass between them, with ESS 0.8262. The equal mixture attains coverage 1, shows less intermodal mass, and reaches ESS 0.9555.

target	KL	$\text{KL}+X_\mu$	$\text{KL}+X_\mu+X_{\hat{\mu}}$	$\text{KL}+X_\mu+X_{(\hat{\mu}+\bar{\nu})/2}$
Two-Moon	0.836 (1.000)	0.911 (1.000)	0.835 (1.000)	0.924 (1.000)
Three-Well	0.988 (0.760)	0.990 (0.760)	0.961 (0.978)	0.989 (0.977)
Himmelblau	0.932 (0.516)	0.981 (0.516)	0.956 (1.000)	0.974 (1.000)
Sparse	0.912 (0.534)	0.927 (0.534)	0.826 (1.000)	0.956 (1.000)

Table 2: Final ESS, with $\kappa = 5$ coverage in parentheses, on the four 2D targets. Bold marks the highest ESS among methods whose sample panel covers every target mode. The scalar coverage diagnostic slightly undercounts the visibly occupied third mode on Three-Well.

The four targets support the intended division of labor. Adding X_μ raises ESS when coverage is already complete. QT samples supply information about modes absent from the annealed batch. Including pushforward samples in the mixture term is associated with cleaner intermodal regions than the $\hat{\mu}$ -only objective. Because there is no $\bar{\nu}$ -only ablation, the last comparison is evidence for the combined mixture rather than a causal estimate of one term. The equal mixture has the highest mode-complete ESS on every target in Table 2.

5.2 High-dimensional product multi-well

We next isolate accuracy scaling over the sixteen dimensions $d = 16, 32, \dots, 256$. Letting $k = \lfloor \log_2 d \rfloor$, we use

$$U(\mathbf{x}) = \frac{1}{2} \|\mathbf{x}\|^2 + 12 \sum_{i=1}^k \exp(-x_i^2). \quad (23)$$

The first k coordinates are symmetric double wells with minima at $\pm \sqrt{\ln 24}$ and barrier approximately 9.9, while the remaining $d - k$ coordinates are standard Gaussian. The target therefore has exactly 2^k equal-weight sign-pattern modes, where $2^k \leq d < 2^{k+1}$. Multimodality occupies only a k -dimensional subspace, but the log-weight error accumulates across all d coordinates.

All four objectives use the same identity-initialized neural spline flow on $[-4, 4]^d$, with 16 bins, 6 transforms, and conditioner widths (256, 256). At every dimension, training uses batch size 250, 2000 Adam steps at learning rate 10^{-3} , gradient clipping at 10^2 , and a fixed validation set of size $5000d$. The live target surrogate uses one AIS update followed by 50 MALA steps of size 2×10^{-3} . This is a standalone one-flow ablation, so its separately reported QT set contains $1000d$ particles rather than the complete validation set used by the staged algorithm. QT uses melt scale 2 and an Armijo L-BFGS quench with initial step 0.5 and 200 iterations. Final validation ESS is evaluated on $\max(200d, 80000)$ fresh source samples. Each objective–dimension pair is one fixed-seed training realization.




Figure 5: ESS on the product multi-well target (23). (a) Final validation ESS over all sixteen tested dimensions. (b) Batch ESS during the $d = 256$ training run; faint curves are raw values and bold curves are 40-step moving averages.

Every expected sign pattern appears in the fresh evaluation samples for every method and dimension, including all 256 patterns at $d = 256$. The ESS comparison is therefore not explained by a missing sign-pattern mode in those samples. Figure 5 shows that every objective with log-ratio variation has higher validation ESS than forward KL at all sixteen dimensions. At $d = 256$, the final values are 0.4322 for forward KL, 0.5608 for $KL+X_\mu$, 0.5979 for $KL+X_\mu+X_{\hat{\mu}}$, and 0.5687 for $KL+X_\mu+X_{(\hat{\mu}+\hat{\nu})/2}$. The two objectives containing $\hat{\mu}$ behave similarly across dimension. During the $d = 256$ training realization, all three objectives with log-ratio variation continue to improve after forward KL begins to plateau.

5.3 Lattice ϕ^4 field theory: ESS at a collective barrier

The tilted scalar ϕ^4 theory tests ESS at a collective barrier. On a periodic $L \times L$ lattice, with $d = L^2$ and magnetization $m(\phi) = d^{-1} \sum_j \phi_j$, the action is

$$U(\phi) = -0.8 \sum_{\langle j,l \rangle} \phi_j \phi_l + \sum_{j=1}^d \left[\phi_j^2 + \frac{1}{2}(\phi_j^2 - 1)^2 \right] + h \sum_{j=1}^d \phi_j.$$

We take $h = 0.0257$ at $L = 6$ and $h = 0.0144$ at $L = 8$. This size-dependent field keeps the extensive tilt scale nearly fixed: $hL^2 = 0.9252$ at $L = 6$ and $hL^2 = 0.9216$ at $L = 8$. Thus the two targets have comparable phase imbalance despite their different volumes. In the broken phase the field orders near two vacua $\pm v\mathbf{1}$, where $v > 0$ is the spontaneous field magnitude and $\mathbf{1} \in \mathbb{R}^d$ is the all-ones vector. A collective domain-wall barrier separates the vacua. The positive linear term gives the two phases unequal probability, so the minority weight $p_+ = \mathbb{P}(m > 0)$ is not fixed by symmetry and must be learned.

Parallel tempering gives the references $p_+ = 0.1251 \pm 0.0009$ and 0.1284 ± 0.0014 , statistically consistent with occupancies 0.1260 and 0.1268 from an independent mirror-MALA construction. That construction samples the positive vacuum of the symmetric $h = 0$ action, obtains the negative vacuum by the exact reflection $\phi \mapsto -\phi$, reweights the pair by the field term, and resamples one million configurations. Each loss

is run at seeds 0, 1, 2 using an identity-initialized neural spline flow on $[-3, 3]^d$, with 16 bins, 6 transforms, and conditioner widths (256, 256). The source is $\mathcal{N}(0, 0.25I_d)$, centered on the barrier. Training uses a fixed validation set of 10^5 source samples, batch size 500, 2000 Adam steps at learning rate 10^{-3} , and gradient clipping at 10^3 . The live surrogate uses one AIS update followed by 50 MALA steps of size 2×10^{-3} . For both objectives containing $X_{\hat{\mu}}$, quench and temper acts on the complete validation set, with melt scale 2 and an Armijo L-BFGS quench using initial step 0.1 and 200 iterations.

	$L = 6$		$L = 8$	
	ESS	p_+	ESS	p_+
forward KL	0.794/0.840/0.826	1/1/0	0.622/0.686/0.684	1/0/1
KL+ X_μ	0.935/0.933/0.945	1/1/0	0.804/0.824/0.821	1/0/0
KL+ $X_\mu+X_{\hat{\mu}}$	0.875/0.873/0.871	0.127/0.128/0.127	0.636/0.645/0.654	0.125/0.124/0.124
KL+ $X_\mu+X_{(\hat{\mu}+\bar{\nu})/2}$	0.896/0.877/0.894	0.127/0.128/0.127	0.690/0.676/0.681	0.131/0.130/0.129
PT reference	—	0.1251 ± 0.0009	—	0.1284 ± 0.0014

Table 3: Final ESS and reweighted minority-phase weight p_+ for three seeds on the two tilted ϕ^4 targets, $(L, h) = (6, 0.0257)$ and $(8, 0.0144)$. The size-dependent field keeps $hL^2 \simeq 0.923$, making the phase imbalance comparable across the two volumes. Integer values 0 and 1 indicate collapse onto one vacuum; bold marks the higher ESS among the objectives that recover both vacua.

Forward KL and KL+ X_μ collapse at every size and seed. Their annealed batches start from pushforward proposals and do not cross the collective barrier within the run, so the second vacuum remains absent and reweighting gives $p_+ = 0$ or 1. Nevertheless, KL+ X_μ has the highest ESS at both sizes: it improves calibration on the reached vacuum while leaving the missing phase undetected by ESS.




Figure 6: Magnetization density on the $L = 8$ lattice at seed 0. Each panel shows the pushforward density (gray), the reweighted density (color), and the resampled mirror-MALA reference density (dotted); the title reports ESS and the reweighted p_+ .

Both objectives using QT populate both vacua at every seed. Their recovered p_+ lies between 0.1238 and 0.1312, within 0.005 of the parallel-tempering estimate and the resampled mirror-MALA reference at the corresponding lattice size. In these finite-capacity fits, recovering both vacua is accompanied by some mass in the suppressed region between phases and lower ESS. At $L = 8$, the equal mixture has ESS 0.6756–0.6902, compared with 0.6355–0.6541 for the $\hat{\mu}$ -only objective. Its samples retain both vacua with less visible bridge mass, consistent with the intended role of the pushforward samples.

5.4 Boltzmann generator on p -state clock model

The preceding benchmarks use a single flow. We now test the full annealed construction on the $p = 6$ clock model on a periodic 8×8 lattice,

$$U(\boldsymbol{\theta}) = - \sum_{\langle jl \rangle} \cos(\theta_j - \theta_l) - \frac{1}{2} \sum_{j=1}^d \cos(6\theta_j), \quad \boldsymbol{\theta} \in [-\pi, \pi)^d, \quad d = 64, \quad (24)$$

where $\langle jl \rangle$ runs over nearest-neighbor pairs. The source is uniform on the torus and the bridge is $U_t = tU$. At low temperature, the target separates into six symmetry-related clock sectors.

Each stage uses an identity-initialized neural circular spline flow with 16 bins, 6 autoregressive transforms, and conditioner widths (256, 256). The validation set contains 4×10^5 particles, and both SMC and quench and temper act directly on this complete set. Every stage uses 3000 Adam steps at learning rate 10^{-3} , a four-level AIS surrogate with 100 MALA steps of size 10^{-3} , energy clipping at 10^3 , and gradient clipping at 10^2 . For KLXX, QT uses wrapped-Gaussian melt scale 2π and an Armijo L-BFGS quench with initial step 10^{-2} and 100 iterations. The clock stage schedule starts from $t_{\text{safe}} = 0.25$, uses enlarge factor $\Gamma = 2$, and applies the nondefault gates $\tau_{\text{SMC}} = 0.7$ and $\tau_v = 0.4$. We report batch sizes $B = 2000, 1000, 500, 250$. An additional completed run lies below this reporting cutoff and is excluded from the table and all aggregate clock claims.

The comparison is schedule-matched. Forward KL first constructs an adaptive stage schedule using SMC preselection and an ESS acceptance gate. The loss $\text{KL} + X_\mu + X_{(\hat{\mu} + \bar{\nu})/2}$ is then trained on the identical accepted values of t , without a second gate. Hence the two methods see the same stage increments within each batch-size pair. Every configuration is a single run, so individual stage ESS values carry single-run fluctuation.

The principal metric is the selected-proposal validation ESS of each accepted stage. The propagation factor in (22) summarizes stagewise weight degeneracy across the stage sequence; smaller is better. It is not a full-chain ESS or an endpoint error estimate. Table 4 reports every accepted stage and the resulting factor.

stage k	1	2	3	4	5	6	7	8	F
t_k ($B = 2000$)	0.250	0.495	0.663	0.779	0.887	1.000	—	—	
forward KL	0.751	0.570	0.471	0.474	0.511	0.602	—	—	33.9
$\text{KL} + X_\mu + X_{(\hat{\mu} + \bar{\nu})/2}$	0.921	0.759	0.608	0.562	0.558	0.649	—	—	11.6
t_k ($B = 1000$)	0.250	0.495	0.613	0.728	0.841	0.952	1.000	—	
forward KL	0.672	0.479	0.580	0.463	0.427	0.520	0.830	—	63.0
$\text{KL} + X_\mu + X_{(\hat{\mu} + \bar{\nu})/2}$	0.866	0.642	0.700	0.545	0.462	0.541	0.859	—	22.0
t_k ($B = 500$)	0.250	0.421	0.590	0.705	0.784	0.895	1.000	—	
forward KL	0.580	0.549	0.408	0.432	0.525	0.417	0.530	—	153.5
$\text{KL} + X_\mu + X_{(\hat{\mu} + \bar{\nu})/2}$	0.778	0.695	0.506	0.500	0.595	0.429	0.566	—	50.7
t_k ($B = 250$)	0.250	0.421	0.539	0.654	0.734	0.811	0.887	1.000	
forward KL	0.468	0.448	0.495	0.428	0.506	0.484	0.514	0.424	421.2
$\text{KL} + X_\mu + X_{(\hat{\mu} + \bar{\nu})/2}$	0.659	0.586	0.601	0.480	0.563	0.536	0.564	0.464	113.8

Table 4: Selected-proposal validation ESS on every accepted stage of the schedule-matched clock-model comparison between forward KL and $\text{KL} + X_\mu + X_{(\hat{\mu} + \bar{\nu})/2}$. Within each batch-size block, both losses traverse the listed stage schedule. The final column is the propagation factor in (22); bold marks the better loss at each stage and the smaller propagation factor.

The loss $\text{KL} + X_\mu + X_{(\hat{\mu} + \bar{\nu})/2}$ raises the geometric mean stage ESS at every reported batch size: from 0.556 to 0.665 at $B = 2000$, from 0.553 to 0.643 at $B = 1000$, from 0.487 to 0.571 at $B = 500$, and from 0.470 to 0.553 at $B = 250$. Correspondingly, it reduces the factor by ratios between 2.87 and 3.70. The gain is not confined to the small-batch regime: at the largest batch, the factor falls from 33.9 to 11.6.

The comparison is also consistent stage by stage. The loss $\text{KL} + X_\mu + X_{(\hat{\mu} + \bar{\nu})/2}$ has higher validation ESS on all 28 shared stages, including every stage in the large-batch runs. As the batch shrinks, the adaptive

forward KL schedule lengthens from six to eight reported stages and both propagation factors grow. KLXX nevertheless accumulates less stagewise weight degeneracy before the final target is reached.




Figure 7: Angle densities from 2×10^6 fresh samples advanced through the trained $B = 2000$ $\text{KL} + \text{X}_\mu + \text{X}_{(\hat{\mu} + \bar{\nu})/2}$ stage sequence by map, reweighting, resampling, and MALA. The upper panel is the single-site marginal. The lower panel gives pair differences at offsets $(0, 1)$, $(1, 1)$, and $(0, 2)$; dots and the inset resolve the one-clock-step angle $\pi/3$. Dashed lines show the uniform density.

The single-site marginal exhibits all six clock wells with nearly equal peak heights. Pair differences are most concentrated at zero for nearest neighbors; as the lattice separation increases, the alignment peak falls and the shoulders near $\pm\pi/3$ gain mass. This is the expected decay of local angular order with distance and confirms that the staged sampler is not confined to a single global sector.

To quantify whether staged resampling preserves the sixfold symmetry, we perform two fresh-rebuild tests through the trained stage sequences. The batch-size test fixes $N = 6.4 \times 10^5$, uses $B \in \{2000, 1000, 500, 250\}$, and averages over four matched sampling seeds for each method and batch size. The particle-count test fixes $B = 2000$ and uses $N = 10^4 2^k$, $k = 0, \dots, 6$, with 2^{8-k} independent rebuilds at each size so every particle-count value has equal total work of 2.56×10^6 particles. For each sample, let s be the clock sector nearest to the argument of the complex magnetization $d^{-1} \sum_{j=1}^d e^{i\theta_j}$, and let p_s be the empirical fraction assigned to sector s . We define

$$\text{occupancy bias} := \frac{1}{6} \sum_{s=0}^5 \left| p_s - \frac{1}{6} \right|.$$




Figure 8: Clock-sector occupancy bias under two fresh-rebuild tests. (a) Mean bias at fixed $N = 6.4 \times 10^5$ across training batch sizes; each bar averages four matched sampling seeds. (b) Mean bias at fixed $B = 2000$ versus particle count; at $N = 10^4 2^k$, each mean uses 2^{8-k} independent rebuilds, keeping the total work per point at 2.56×10^6 particles. Error bars in both panels are ± 2 standard errors. The dashed line in (b) is an $N^{-1/2}$ reference anchored at the smallest $\text{KL} + X_\mu + X_{(\hat{\mu} + \hat{\nu})/2}$ point. The panels use separately generated rebuild ensembles.

At fixed N , panel (a) shows that the mean occupancy bias increases as the training batch shrinks for both methods. KLXX has the lower mean at $B = 2000$, 1000 , and 250 . At those batch sizes, the forward KL means are 0.00573 , 0.00888 , and 0.02760 , while the KLXX means are 0.00235 , 0.00517 , and 0.01506 . At $B = 500$ the ordering reverses slightly, from 0.00963 for forward KL to 0.01082 for KLXX, and the intervals defined by ± 2 standard errors overlap.

At fixed $B = 2000$, panel (b) follows the Monte Carlo rate: the log-log slopes of occupancy bias against N are -0.459 for forward KL and -0.479 for $\text{KL} + X_\mu + X_{(\hat{\mu} + \hat{\nu})/2}$. The absence of a visible bias floor is consistent with sampling fluctuation rather than a particle-count-independent sector imbalance over the tested range. The KLXX curve is lower at all seven reported sizes, reducing the mean bias from 0.03735 to 0.02397 at $N = 10^4$ and from 0.00552 to 0.00358 at $N = 6.4 \times 10^5$. Every reported rebuild in both panels retains all six sectors. Together with the six peaks in Figure 7, these tests show that the staged samples are not confined to one global sector. The rebuild seeds measure sampling variability; each stage schedule for a given method and batch size is one training run, so the between-method differences do not estimate variation over independently trained models.

6 Boltzmann generators on molecular potential energies

The model problems above separate the target and QT training signals on explicit targets and probe the term built from pushforward samples. Molecular systems add a singular nonbonded core, a mixed Euclidean-periodic domain, and more conformational degrees of freedom as molecular size grows. We first define the molecular measure and its regularized potential, then add sharpening to the staged generator and report the completed alkane family. The pending targets are classified below as achiral or as chiral with one fixed stereochemical configuration.

6.1 Molecular potential models

For Cartesian coordinates x , let $E(x)$ denote the complete molecular-mechanics energy

$$E(x) = E_{\text{bond}} + E_{\text{angle}} + E_{\text{torsion}} + E_{\text{LJ}} + E_{\text{Coul}} + E_{\text{GB}}^{\text{OBC1}} + E_{\text{ACE}}. \quad (25)$$

The bonded and vacuum nonbonded part has the Amber form (Cornell et al., 1995)

$$E_{\text{vac}}(x) = \sum_b k_b (r_b - r_b^0)^2 + \sum_a k_a (\theta_a - \theta_a^0)^2 + \sum_\tau \frac{V_\tau}{2} [1 + \cos(n_\tau \varphi_\tau - \gamma_\tau)] \\ + \sum_{\substack{i < j \\ \text{nonbonded}}} \left\{ 4\epsilon_{ij} \left[\left(\frac{\sigma_{ij}}{r_{ij}} \right)^{12} - \left(\frac{\sigma_{ij}}{r_{ij}} \right)^6 \right] + k_e \frac{q_i q_j}{r_{ij}} \right\}, \quad (26)$$

Here b , a , and τ index bonds, angles, and proper or improper torsions, while i, j index nonbonded atom pairs after topology-defined exclusions and scaling. The symbols r_b^0 , θ_a^0 , k_b , k_a , V_τ , n_τ , and γ_τ are force-field equilibrium values and coefficients; σ_{ij} , ϵ_{ij} , and q_i are the nonbonded parameters. OBC1 supplies the generalized-Born electrostatic solvation term (Onufriev et al., 2004), and ACE supplies the nonpolar term (Schaefer and Karplus, 1996). All terms in (25) belong to the sampled potential.

A nonlinear molecule with A atoms has $d = 3A - 6$ internal degrees of freedom after rigid translation and rotation are removed. A fixed internal-coordinate tree reconstructs $x = x(q)$ from whitened bond and angle coordinates and periodic torsions. Write d_E and d_T for the numbers of nonperiodic Euclidean coordinates and periodic torsions, respectively, and let $\mathbb{T} = \mathbb{R}/(2\pi\mathbb{Z})$. Consequently,

$$\mathcal{Q} = \mathbb{R}^{d_E} \times \mathbb{T}^{d_T}, \quad d_E + d_T = 3A - 6, \quad (27)$$

and the physical reduced potential at inverse temperature $\beta_T = (k_B T)^{-1}$ is

$$U(q) = \beta_T E(x(q)) - \log J(q), \quad (28)$$

where $J(q)$ is the Cartesian volume Jacobian of the reconstruction, including the whitening maps. The source is Gaussian on the Euclidean block and uniform on the torus. The mixed spline flow uses ordinary rational-quadratic splines on the Euclidean coordinates and circular splines on the torsions.

target	atoms / d	model	status
<i>Achiral targets</i>			
alkanes, C ₁ –C ₆	5–20/9–54	GAFF2/AM1-BCC; OBC1/ACE	complete
<i>N</i> -methylacetamide (NMA)	12/30	ff96; OBC1/ACE	pending
glycerol	14/36	GAFF2/AM1-BCC; OBC1/ACE	pending
neutral diethanolamine	18/48	GAFF2/AM1-BCC; OBC1/ACE	pending
<i>Chiral targets: one fixed stereochemical configuration</i>			
(<i>S</i>)-2-butanol	15/39	GAFF2/AM1-BCC; OBC1/ACE	pending
(2 <i>R</i> , 3 <i>R</i>)-2,3-butanediol	16/42	GAFF2/AM1-BCC; OBC1/ACE	pending
alanine dipeptide	22/60	ff96; OBC1/ACE	pending

Table 5: Molecular targets in mixed internal coordinates, grouped by configurational stereochemistry. All systems are at 300 K, use no nonbonded cutoff or bond constraints, and include the coordinate Jacobian in (28). For each chiral target, the generator and reference are restricted to one specified stereochemical configuration, while all conformations within that configuration remain part of the target. Alkane observables use a two-seed OpenMM reference; the other generators remain pending.

The alkane family uses neutral all-atom bundles with GAFF2 parameters, AM1-BCC charges, graph-symmetry and neutrality projection, and mbondi2 radii. GAFF and AM1-BCC are described in Wang et al. (2004) and Jakalian et al. (2002); the bundles are prepared with AmberTools (Case et al., 2023). OBC1/ACE is used across the homologous series, and the independent references are propagated with OpenMM (Eastman et al., 2017). NMA uses the 12-atom ACE–NME ff96/OBC1/ACE Hamiltonian. The two chiral alcohols use GAFF2/AM1-BCC with OBC1/ACE. Alanine dipeptide uses the 22-atom ACE–ALA–NME Hamiltonian of the FAB benchmark (Midgley et al., 2022): Amber ff96, OBC1 with igb = 2, mbondi2 radii, solute and solvent dielectric constants 1 and 78.5, and the L-alanine component selected by the reference trajectory.

The alkanes, glycerol, neutral diethanolamine, and NMA have no fixed configurational stereocenter under these bundle definitions. Neutral diethanolamine retains both signs of its nitrogen-pyramidal coordinate, and the NMA amide torsion remains periodic so that both cis and trans conformations belong to its target. The chiral targets retain only the named configurations: S for 2-butanol, $(2R, 3R)$ for 2,3-butanediol, and L-alanine for alanine dipeptide. All conformational basins within each selected configuration remain part of the corresponding target. These support conventions are part of the target rather than an after-the-fact filter.

The Lennard–Jones core in (26) grows as r^{-12} , so early flow proposals can place a few samples at energies and gradients far beyond the scale of the thermally occupied region. We therefore use a two-parameter molecular surrogate. Let x_* be the bundle reference geometry. After flooring only the regular and exception Coulomb/Lennard–Jones pair distances at r , denote the complete force-field energy by $E_r(x)$, its reference value by $E_r^* := E_r(x_*)$, and its reference-relative value by $\Delta E_r(x) := E_r(x) - E_r^*$. For an energy threshold e , define the soft cap

$$\mathcal{C}_e(\Delta E) = \begin{cases} \Delta E, & \Delta E \leq e, \\ e[1 + \log(\Delta E/e)], & \Delta E > e, \end{cases} \quad (29)$$

and collect the two regularization parameters in $\rho = (e, r)$. The corresponding reduced potential is

$$U^{\{\rho\}}(q) = \beta_T [E_r^* + \mathcal{C}_e(\Delta E_r(x(q)))] - \log J(q). \quad (30)$$

The threshold e is measured in kJ mol^{-1} , and the pair floor r is measured in nm. The component formulas of the bonded and OBC1/ACE terms are unchanged before the scalar cap is applied to the complete energy. The coordinate Jacobian and domain are not capped. The physical potential U in (28) is recovered in the formal limit $e \rightarrow \infty$, $r \rightarrow 0$. Thus $U^{\{\rho\}}$ is a sampling bridge, not optimizer clipping or a replacement for the physical energy used in the observables.

6.2 Boltzmann generators with sharpening

A softer molecular potential can make early stages easier in a given adaptive-staging run. After each flow stage, *sharpening* reweights the complete validation set toward a stricter regularization setting, resamples it, and applies MALA under the stricter potential. This is a stochastic potential update, not an additional flow-training loss. To extend the source-to-target interpolation of Section 4, let $\rho_0 = (e_0, r_0)$ and $\rho_1 = (e_1, r_1)$ be the soft and final settings and interpolate them by

$$\rho_s = (1 - s)\rho_0 + s\rho_1, \quad s \in [0, 1]. \quad (31)$$

The molecular experiments use the finite endpoint $\rho_1 = (100, 0.15)$ and audit its residual difference from U directly. A fixed-regularization run has $\rho_0 = \rho_1$. Extending the Section 4 bridge by the single index s , set

$$U_{t,s}(q) = (1 - t)U_0(q) + tU^{\{\rho_s\}}(q), \quad \mu_{t,s} \propto \exp(-U_{t,s}). \quad (32)$$

Thus t remains the source-to-target interpolation parameter defined in Section 4, while s changes only the regularization setting. If s is fixed, suppressing the second index recovers the notation U_t, μ_t of Section 4.

Accepted sample sets lie on the diagonal $s = t$, but one stage advances the two coordinates in sequence:

$$\mathcal{Y}_{k-1} \approx \mu_{t_{k-1}, t_{k-1}} \xrightarrow[\text{selected } (\tilde{\mathcal{Y}}_k, w_k)]{\text{flow or identity}} \mathcal{Y}_k^- \approx \mu_{t_k, t_{k-1}} \xrightarrow[w_k^{\text{sharp}}]{\text{sharpen}} \mathcal{Y}_k \approx \mu_{t_k, t_k}. \quad (33)$$

The first arrow holds $s = t_{k-1}$ fixed and advances t from t_{k-1} to t_k . Potential-space SMC and mixed-domain MALA carry a copy of all N particles along this bridge. The resulting sample set is held fixed and subsampled during flow training; its smallest incremental ESS also screens the proposed t_k . In KLXX, a second sample set is built by applying QT under $U_{t_k, t_{k-1}}$ to N fresh source draws. Thus the molecular implementation uses a fixed SMC target set and a fresh-source QT set, rather than the per-stage target surrogate and validation-seeded QT in Section 4.

After training, the candidate flow and identity proposals are evaluated on the complete carried set $\mathcal{Y}_{k-1}$. The pair with larger ESS over the complete validation set is denoted by $(\tilde{\mathcal{Y}}_k, w_k)$; ties retain the identity.

Reweighting, resampling, and MALA under $U_{t_k, t_{k-1}}$ produce $\mathcal{Y}_k^-$. The second arrow holds $t = t_k$ fixed and advances the regularization coordinate from t_{k-1} to t_k . Its exact log weight is

$$\log w_k^{\text{sharp}}(q) = U_{t_k, t_{k-1}}(q) - U_{t_k, t_k}(q) = t_k \left[U^{\{\rho_{t_{k-1}}\}}(q) - U^{\{\rho_{t_k}\}}(q) \right]. \quad (34)$$

Sharpening is accepted only when its ESS over the complete validation set clears the stage threshold. The samples are then resampled and rejuvenated by MALA under U_{t_k, t_k} . A failed proposal or sharpening gate shrinks $t_k - t_{k-1}$ and retries. The selected map advances the source-to-target bridge; the stochastic sharpening step changes the molecular potential.

We compare three generators. ID performs both stochastic arrows without training a flow. KL trains the forward KL and then selects the trained or identity proposal. KLXX uses the same selection rule with log-ratio variation. All three methods retain the same SMC selection, reweighting, resampling, and MALA steps. Consistently with (22), the molecular propagation factor is

$$F = \prod_{k=1}^K \text{ESS}(w_k)^{-1} \text{ESS}\left(w_k^{\text{sharp}}\right)^{-1}, \quad (35)$$

where $w_k^{\text{sharp}} \equiv 1$ on a fixed regularization path. As in (22), F summarizes incremental weight degeneracy; it is not a fresh source-to-endpoint ESS or an endpoint error estimator.

6.3 Alkane family results

We test methane through *n*-hexane at 300 K, with mixed-coordinate dimensions 9, 18, $\dots$, 54. Every configuration uses one seed and reaches $t = 1$. All endpoint sample sets target $(e, r) = (100, 0.15)$. Direct runs start and remain at this endpoint. Sharpening runs begin at (50, 0.15) for propane, (50, 0.20) for butane and pentane, and (50, 0.25) for hexane, then advance linearly to the common endpoint. The three methods use method-specific adaptive stage schedules, so stage counts and recorded wall times are descriptive rather than fixed-compute comparisons. The reported molecular wall times were recorded on a single NVIDIA GeForce RTX 5090 GPU.

All molecular flows use six mixed spline transforms, 32 bins, conditioner widths (256, 256), spline bound 8, and minimum slope 10^{-3} , and are initialized at the identity. The validation and batch sizes (N, B) are (100000, 5000) for methane, (120000, 6000) for ethane, (160000, 8000) for propane, and (200000, 10000) for butane through hexane. Each accepted flow stage uses 500 Adam steps for methane through propane, 400 for butane, and 300 for pentane and hexane, with learning rate 10^{-3} and 25 warmup steps. Potential-space SMC uses eight intermediate levels; mixed-domain MALA uses 100 steps of size 10^{-3} . QT uses melt scale 1, 200 L-BFGS steps with initial step 10^{-2} , and the complete molecular QT set. The safe start is 0.25 for methane and ethane and 0.1 for the larger molecules. The validation gate is 0.4 throughout; the SMC gate is 0.3 for methane and ethane, 0.5 for propane, and 0.8 for butane through hexane.

molecule (d)	start ρ_0	ID			KL			KLXX		
		F	stages	time	F	stages	time	F	stages	time
methane (9)	(100, 0.15)	14.69	8	1.4	1.416	3	1.2	1.053	3	2.4
ethane (18)	(100, 0.15)	62.98	9	3.5	6.349	3	1.8	1.423	3	4.1
propane (27)	(100, 0.15)	182.8	11	7.7	181.8	11	13.4	7.738	7	23.0
	(50, 0.15)	95.14	10	6.9	95.02	10	12.1	7.778	6	18.3
butane (36)	(100, 0.15)	367.3	12	13.3	367.1	12	20.3	28.06	10	49.0
	(50, 0.20)	226.4	10	10.8	224.8	10	14.9	14.91	10	46.1
pentane (45)	(100, 0.15)	733.2	13	19.2	754.7	13	24.6	82.71	12	60.4
	(50, 0.20)	363.6	10	13.9	370.2	10	17.2	48.83	9	38.5
hexane (54)	(100, 0.15)	2043	13	23.3	2077	13	29.0	1026	12	69.9
	(50, 0.25)	910.5	10	17.4	891.5	10	22.7	201.5	10	60.1

Table 6: Alkane Boltzmann generator training results, including propagation factor F from (35), the number of accepted stages, and total wall time in minutes. ID applies the identity map at every stage, whereas KL and KLXX may select a trained flow when it improves validation ESS. Each row reports one seed and one regularization path; bold marks the lowest F among the three methods for that path. Every path ends at (100, 0.15), and rows that start at (100, 0.15) use no sharpening.

Within the single-seed runs, KLXX has the smallest factor in all ten matched settings. KL selects trained maps through methane and part of ethane. From propane onward, every accepted KL stage selects the identity proposal, and its factor is close to the ID result. KLXX selects the trained map at every accepted stage through pentane. At hexane, it selects trained maps in 6 of 12 fixed-path stages and 8 of 10 sharpening stages.

The difference between the fixed and sharpening paths increases with molecule size in these runs. The propane KLXX factors are nearly equal, whereas the fixed-path factor divided by the sharpening-path factor is 1.88, 1.69, and 5.09 for butane, pentane, and hexane. The corresponding KLXX sharpening runs take 3.10, 2.24, and 2.65 times the recorded KL wall time. Thus the KLXX sharpening paths have the lowest factor in the larger-molecule comparisons, with additional QT, reweighting, and flow-training cost.

The runs stop at the finite regularization ρ_1 rather than at the formal physical limit. For endpoint samples $q_1, \dots, q_N$, define the endpoint-to-physical weights and their normalized ESS, called the regularization ESS (RESS), by

$$w_i^{\text{phys}} := \exp \left[U^{\{\rho_1\}}(q_i) - U(q_i) \right], \quad \text{RESS} := \text{ESS}(w^{\text{phys}}). \quad (36)$$

The endpoint-to-physical weights are nearly uniform on the final KLXX sample sets: $1 - \text{RESS} = 0$ numerically from methane through butane, 9.5×10^{-6} for pentane, and 2.3×10^{-5} for hexane. These values measure overlap on the sampled sets; they do not prove global equality between $U^{\{\rho_1\}}$ and U .

The KLXX curves are unweighted summaries of the finite-endpoint sample sets; the physical energy E is evaluated only as an observable, without endpoint-to-physical resampling. The physical-potential reference uses two seeds and a 0.25 fs timestep. Methane through propane use 300 K Langevin trajectories; butane through hexane use 300–800 K replica exchange to mix the backbone rotamers. Across the six molecules, the maximum relative difference between KLXX and reference physical-energy means is 0.644%. The largest relative difference among the nonzero carbon radius-of-gyration and terminal-distance means is 1.058%, and the largest absolute rotamer-fraction difference is 0.051. These results show agreement for the selected energy, size, and rotamer summaries across the C_1 – C_6 series. They do not establish equality of the complete distributions.

All thirty ID, KL, and KLXX run manifests used in Table 6 are complete, but each configuration has one training seed. The reference frames are correlated, and the two-seed ranges are descriptive rather than framewise confidence intervals. The table therefore compares the recorded runs; it does not estimate run-to-run variation or distance between the complete equilibrium measures.




Figure 9: Macroscopic observables from the smallest-factor KLXX sample set for each alkane and an independent physical-potential reference. Fixed-path KLXX is used for methane through propane and sharpening KLXX for butane through hexane. (a) Physical energy evaluated under the exact, unregularized potential. (b) Carbon-skeleton radius of gyration $R_{g,C}$ and terminal-carbon distance $R_{ee,C}$. (c) Pooled backbone trans and gauche fractions. Filled solid curves and unhatched bars denote KLXX; open dashed curves and hatched bars denote the reference.

6.4 Achiral molecules: pending

The pending achiral targets, ordered by dimension, are NMA, glycerol, and neutral diethanolamine. NMA uses ff96/OBC1/ACE and has dimension 30. Glycerol and neutral diethanolamine have dimensions 36 and 48, respectively. Both use the alkane construction, with GAFF2/AM1-BCC parameters and OBC1/ACE solvent. A production comparison must resolve the energy distribution and the relevant joint torsional structure rather than only one-dimensional marginals. In particular, NMA must report its amide cis/trans and planarity statistics, while glycerol and neutral diethanolamine require joint rotamer and unrestricted signed-volume diagnostics. These generator runs and matched physical-potential references remain pending and are not included in the claims above.

6.5 Chiral molecules: pending

The smaller chiral targets are (*S*)-2-butanol at dimension 39 and (2*R*,3*R*)-2,3-butanediol at dimension 42. Both use GAFF2/AM1-BCC with OBC1/ACE. Their target supports retain only the named stereochemical configuration; the butanol target fixes one center, while the butanediol target fixes both centers jointly. The planned comparisons must report signed-volume acceptance together with the energy and torsional distributions.

Alanine dipeptide provides the higher-dimensional chiral benchmark with a published reference. Its mixed domain is $\mathbb{R}^{42} \times \mathbb{T}^{18}$, restricted to the L-alanine component represented by the FAB million-frame trajectory. Thus only one stereochemical configuration is sampled, but its Ramachandran and remaining torsional degrees of freedom vary within that component. The intended comparison uses the exact ff96/OBC1/ACE Hamiltonian and reports the Ramachandran (ϕ, ψ) distribution together with full-energy reweighting and the remaining torsions. All three chiral generators remain pending, so none of their sampling results enters this revision.

7 Conclusions

Forward KL is mass-covering for exact distributions, but its sampled loss sees only the target regions supplied to it. This gap can produce a misleading ESS, transfer surrogate error to the flow, and leave target-poor proposal regions weakly constrained. Log-ratio variation addresses these failures by comparing the exact log-density ratio at pairs of samples. The sampling weight determines which regions are compared: target samples measure density variation on reached target regions, QT samples add candidate modes, and pushforward samples expose regions occupied by the proposal. KLXX combines these comparisons while

retaining the exact target as a minimizer and the flow as a tractable importance proposal.

The theory establishes a narrower statement. For the selected Fisher–Rao subgradient flow of forward KL with X_μ , the regularizer adds a nonpositive term to the instantaneous KL dissipation. For a fixed smooth surrogate and $\lambda < \frac{1}{2}$, it also improves the stationary-error bound by a nonnegative discrepancy term. This does not prove finite-step behavior for full KLXX, Adam, or the experimental value $\lambda = 1$.

The numerical results provide the remaining evidence. The 2D and ϕ^4 ablations are consistent with the roles of the target, QT, and pushforward samples; the product test shows higher validation ESS for every tested objective with log-ratio variation across the tested dimensions. On the shared clock stage schedules, KLXX has higher validation ESS on all 28 reported stages and a smaller propagation factor at every reported batch size. In the single-seed alkane runs, KLXX has the smallest factor in every matched setting, while forward KL approaches the identity baseline as molecule size grows. The sharpening paths have lower KLXX factors for the larger alkanes. Selected energy, size, and rotamer summaries agree with independent physical-potential simulations. These findings are consistent with an explanation based on additional support information; they are not a theorem of universal dominance or a fixed-compute comparison.

The clock and alkane configurations require additional training seeds to measure run-to-run variation. The finite molecular endpoint has nearly uniform endpoint-to-physical weights on the reported sample sets, but a continuation to U would remove the remaining regularization. The molecular targets marked pending in Table 5 still require generator runs and matched references; the alanine study will also permit a direct comparison with FAB (Midgley et al., 2022). Larger molecules and explicit solvent remain outside the present results.

Acknowledgments

The normalizing flow computations use `jflows` and `jflows_md`. Molecular benchmark data were generated with OpenMM and AmberTools, and the alanine dipeptide reference data were taken from the FAB study (Midgley et al., 2022). The experiment sources and software repositories are available at:

```
https://github.com/xuda-ye-math/X-regularization
https://github.com/xuda-ye-math/jflows
https://github.com/xuda-ye-math/jflows_md
```

The GPU-accelerated runs used one NVIDIA GeForce RTX 5090 with 32 GB of VRAM.

A Fisher–Rao gradient flow with log-ratio variation

Section 2 states the main gradient-flow results used in the paper. This appendix supplies their derivations and proofs on the space of densities. We view ν as a generic probability density evolving along the gradient flow of the loss under an appropriate Riemannian metric on the space of densities (Ambrosio et al., 2008; Villani, 2009) and ask how fast ν_t converges to the target μ in $\text{KL}(\mu\|\nu_t)$.

The absolute value in X_μ is not differentiable when two log-ratios coincide. We set $\text{sgn}(0) = 0$, selecting the zero subgradient at a tie, and study the corresponding formal Fisher–Rao subgradient flow. The densities below are assumed positive and regular enough for the displayed calculations.

In this section we restrict to the loss with only the forward KL and the shape regularizer X_μ , the log-ratio variation at weight $\omega = \mu$,

$$\mathcal{L}[\nu] = \text{KL}(\mu\|\nu) + \lambda X_\mu(\mu\|\nu) = \int_{\mathbb{R}^d} \mu(x) \log \frac{\mu(x)}{\nu(x)} dx + \lambda \iint_{\mathbb{R}^d \times \mathbb{R}^d} \mu(x) \mu(y) \left| \log \frac{\mu(x)}{\nu(x)} - \log \frac{\mu(y)}{\nu(y)} \right| dx dy, \quad (37)$$

where $\lambda > 0$ is fixed. The calculations below hold along a positive solution of the selected flow; when $\lambda \leq \frac{1}{2}$, the birth term in (43) is nonnegative and preserves positivity. Throughout this section we abbreviate the importance weight by

$$w(x) := \frac{\mu(x)}{\nu(x)},$$

so that the log-ratio is $z(x) = \log w(x) = \log \mu(x) - \log \nu(x)$, the field z of (6). The importance weight w is distinct from the localization measure ω in X_ω .

The choice of Riemannian metric matters. In the Wasserstein-2 geometry, induced by optimal transport (Jordan et al., 1998; Otto, 2001), the forward KL fails to be displacement convex along Wasserstein-2 geodesics, and the standard Otto–Villani/Bakry–Émery argument (Otto and Villani, 2000; Bakry et al., 2014) used for the reverse KL does not apply. We therefore work in the Fisher–Rao (FR) geometry (Amari, 2016; Bauer et al., 2016). Along the selected positive flow, the forward KL has an exponential bound and X_μ contributes a nonlocal birth–death term (Lu et al., 2019): the update at x depends on how its importance weight ranks among weights under μ .

Denoting the ν -mean by $\langle \cdot \rangle_\nu := \int (\cdot) \nu dx$, the Fisher–Rao gradient of a functional $\mathcal{F}[\nu]$ and the associated dissipation identity along the FR gradient flow $\partial_t \nu_t = -\text{grad}^{\text{FR}} \mathcal{F}(\nu_t)$ are

$$\text{grad}^{\text{FR}} \mathcal{F}(\nu) = \nu \left(\frac{\delta \mathcal{F}}{\delta \nu} - \left\langle \frac{\delta \mathcal{F}}{\delta \nu} \right\rangle_\nu \right), \quad \frac{d}{dt} \mathcal{F}[\nu_t] = - \int_{\mathbb{R}^d} \left(\frac{\delta \mathcal{F}}{\delta \nu} - \left\langle \frac{\delta \mathcal{F}}{\delta \nu} \right\rangle_{\nu_t} \right)^2 \nu_t dx. \quad (38)$$

The mean-subtraction enforces $\int \text{grad}^{\text{FR}} \mathcal{F} = 0$, so the FR flow conserves probability mass.

A.1 Derivation of X_μ -regularized Fisher–Rao gradient flow

The first variation of the forward KL with respect to ν is

$$\frac{\delta \text{KL}(\mu\|\nu)}{\delta \nu}(x) = -\frac{\mu(x)}{\nu(x)} = -w(x), \quad (39)$$

so by (38), using $\langle w \rangle_\nu = \int_{\mathbb{R}^d} \mu dx = 1$, the FR gradient of the forward KL is

$$\text{grad}^{\text{FR}} [\text{KL}(\mu\|\nu)](x) = \nu(x) - \mu(x). \quad (40)$$

For X_μ , the selected subvariation applies $\text{sgn}(0) = 0$ to the absolute value. At an evaluation point r , the variation of the inner log-ratio difference is

$$\frac{\delta}{\delta \nu(r)} [\log w(x) - \log w(y)] = -\frac{\delta(r-x)}{\nu(x)} + \frac{\delta(r-y)}{\nu(y)}.$$

Here δ is the Dirac distribution, so integration against $\delta(r - x)$ evaluates the remaining integrand at $x = r$. Applying the selected sign subgradient and then exchanging x and y in the second term gives

$$\begin{aligned}\frac{\delta X_\mu(\mu\|\nu)}{\delta\nu}(r) &= \iint_{\mathbb{R}^d \times \mathbb{R}^d} \mu(x)\mu(y) \operatorname{sgn}(w(x) - w(y)) \left[-\frac{\delta(r-x)}{\nu(x)} + \frac{\delta(r-y)}{\nu(y)} \right] dx dy \\ &= -2w(r) \int_{\mathbb{R}^d} \mu(y) \operatorname{sgn}(w(r) - w(y)) dy.\end{aligned}$$

Hence the selected subvariation is

$$\frac{\delta X_\mu(\mu\|\nu)}{\delta\nu}(x) = -2w(x)S(x), \quad S(x) := \mathbb{E}_{y \sim \mu} [\operatorname{sgn}(w(x) - w(y))]. \quad (41)$$

We call S the *rank field* of the importance weight under μ : $S(x) > 0$ indicates that $w(x)$ is comparatively large within the distribution of weights, and hence that ν pointwise undercovers μ relative to that distribution; $S(x) < 0$ has the opposite interpretation, and $S \equiv 0$ at $\nu = \mu$. By antisymmetry,

$$\mathbb{E}_\mu[S] = \mathbb{E}_{x, y \sim \mu} [\operatorname{sgn}(w(x) - w(y))] = 0,$$

and the selected FR subgradient of X_μ with respect to ν is

$$\operatorname{grad}^{\text{FR}}[X_\mu(\mu\|\nu)](x) = -2\mu(x)S(x). \quad (42)$$

The selected FR subgradient of $\mathcal{L}[\nu] = \text{KL}(\mu\|\nu) + \lambda X_\mu(\mu\|\nu)$ is the sum of (40) and (42):

$$\operatorname{grad}^{\text{FR}}[\mathcal{L}](x) = \nu(x) - \mu(x) - 2\lambda\mu(x)S(x).$$

The selected Fisher–Rao subgradient flow therefore reads

$$\partial_t \nu_t(x) = \mu(x) - \nu_t(x) + 2\lambda\mu(x)S_t(x), \quad \text{where } S_t(x) := \mathbb{E}_{y \sim \mu} \left[\operatorname{sgn} \left(\frac{\mu(x)}{\nu_t(x)} - \frac{\mu(y)}{\nu_t(y)} \right) \right]. \quad (43)$$

Equation (43) is a *nonlocal* PDE: its right-hand side at x depends on $S_t(x)$, which in turn requires the entire profile of the importance weight μ/ν_t over the support of μ . Mass is conserved along the flow because

$$\frac{d}{dt} \int_{\mathbb{R}^d} \nu_t(x) dx = 2\lambda \int_{\mathbb{R}^d} \mu(x) S_t(x) dx = 0,$$

and the flow is stationary at $\nu_t = \mu$, where $\mu(x)/\nu_t(x) \equiv 1$ and hence $S_t \equiv 0$. When $\lambda = 0$, (43) reduces to the pure Fisher–Rao gradient flow of the forward KL.

A.2 Convergence analysis of the X_μ -regularized flow

To analyze the convergence rate of the regularized flow (43), we take the forward KL $\text{KL}(\mu\|\nu_t)$ as the error metric and differentiate it along the flow:

$$\frac{d}{dt} \text{KL}(\mu\|\nu_t) = - \int_{\mathbb{R}^d} \frac{\mu^2}{\nu_t} (1 + 2\lambda S_t) dx + \int_{\mathbb{R}^d} \mu dx = -\chi^2(\mu\|\nu_t) - 2\lambda \int_{\mathbb{R}^d} \frac{\mu^2}{\nu_t} S_t dx,$$

where the chi-squared divergence is

$$\chi^2(\mu\|\nu) := \int_{\mathbb{R}^d} (w - 1)^2 \nu dx = \int_{\mathbb{R}^d} w^2 \nu dx - 1.$$

The sign of the second integral is settled by the following rank-correlation identity.

Lemma 1 *Let μ and ν be positive probability densities for which the displayed integral is finite. Then*

$$\int_{\mathbb{R}^d} \frac{\mu(x)^2}{\nu(x)} S(x) dx = \frac{1}{2} \mathbb{E}_{x, y \sim \mu} [|w(x) - w(y)|] \geq 0,$$

with equality if and only if w is μ -almost-everywhere constant.

Proof. Since $\mu^2/\nu = \mu w$, the definition of S gives

$$\int_{\mathbb{R}^d} \frac{\mu^2}{\nu} S dx = \mathbb{E}_\mu[w(x)S(x)] = \mathbb{E}_{x,y \sim \mu}[w(x) \operatorname{sgn}(w(x) - w(y))].$$

Symmetrizing in $x \leftrightarrow y$ gives

$$\int_{\mathbb{R}^d} \frac{\mu^2}{\nu} S dx = \frac{1}{2} \mathbb{E}_{x,y \sim \mu}[(w(x) - w(y)) \operatorname{sgn}(w(x) - w(y))] = \frac{1}{2} \mathbb{E}_{x,y \sim \mu}[|w(x) - w(y)|].$$

□

Applying Lemma 1 with $\nu = \nu_t$ and $S = S_t$, the dissipation rewrites as

$$\frac{d}{dt} \operatorname{KL}(\mu \| \nu_t) = -\chi^2(\mu \| \nu_t) - \lambda \mathbb{E}_{x,y \sim \mu} \left[\left| \frac{\mu(x)}{\nu_t(x)} - \frac{\mu(y)}{\nu_t(y)} \right| \right]. \quad (44)$$

The elementary inequality $w \log w \leq w^2 - w$ applied to $w = \mu/\nu_t$ gives $\chi^2(\mu \| \nu_t) \geq \operatorname{KL}(\mu \| \nu_t)$, hence $\frac{d}{dt} \operatorname{KL}(\mu \| \nu_t) \leq -\operatorname{KL}(\mu \| \nu_t)$, and Grönwall's lemma yields the exponential decay

$$\operatorname{KL}(\mu \| \nu_t) \leq e^{-t} \operatorname{KL}(\mu \| \nu_0), \quad t \geq 0, \quad (45)$$

with no structural assumption on μ : neither log-concavity, nor a spectral gap, nor any growth condition. This contrasts with the Wasserstein-2 geometry, where the forward KL is not displacement convex and an analogous exponential rate is unavailable without further assumptions.

The uniform convergence rate (45) should *not* be read as a blanket statement that minimizing the forward KL is intrinsically faster or more efficient than minimizing the reverse KL. The FR flow is a continuous-time idealization on the space of densities and does not account for sampling cost. In practice, evaluating the forward KL and its gradient requires samples from μ , and producing representative samples in high dimensions is itself difficult. The empirical KLXX objective therefore also evaluates the exact log-ratio on the quench and temper measure $\hat{\mu}$, which is intended to supply support information that may be absent from the target-sample measure. The bound above quantifies convergence given exact target expectations; it does not establish the behavior of this additional empirical term or remove the cost of constructing its samples.

A naive comparison of dissipation rates between the unregularized KL loss and $\operatorname{KL} + \lambda X_\mu$ would also be misleading, since scaling the loss by a constant is mere time reparametrization with no qualitative gain. The substantive contribution of X_μ is structural rather than rate-based: the pure forward KL term depends locally on the value of $\mu(x)/\nu_t(x)$, whereas the rank field S_t compares that value with the complete distribution of importance weights under μ . The regularized flow (43) therefore couples the update at x to global rank information that the local forward KL term does not contain.

Two further cautions concern the discrete realization of the flow. First, the time-reparametrization argument cuts in both directions once the optimizer enters. Training follows Adam steps, not the FR flow: the per-coordinate second-moment normalization absorbs a global rescaling of the loss, so a larger gradient buys no speedup by itself, while a regularizer that reshapes the relative magnitudes of the gradient components does change the optimizer's effective time scale. Part of any observed difference in training speed between KL and $\operatorname{KL} + \lambda X_\mu$ may therefore be optimizer scaling rather than information geometry, and we claim no universal transient speedup on this basis. Second, the dissipation (44) is a density-space statement, while each gradient step uses a finite, generally biased sample estimate. The staged clock experiment examines this finite-batch setting with a schedule-matched comparison: its quantitative claims use the selected-proposal ESS over the complete validation set at each stage, not a composed endpoint estimate or a diagnostic computed on a training mini-batch.

Remark 4 (Monotone decay of X_μ) *The regularizer X_μ also decreases monotonically along the selected regularized flow: pairing the selected subvariation $\delta X_\mu / \delta \nu = -2wS$ (41) against the flow (43) gives*

$$\frac{d}{dt} X_\mu(\mu \| \nu_t) = -2 \int_{\mathbb{R}^d} \frac{\mu^2}{\nu_t} S_t dx - 4\lambda \int_{\mathbb{R}^d} \frac{\mu^2}{\nu_t} S_t^2 dx \leq 0.$$

The first term is nonpositive by Lemma 1 and vanishes only when μ/ν_t is μ -almost-everywhere constant. Hence the inequality is strict away from that case, and X_μ is driven to zero along the flow.

A.3 Asymptotic accuracy for a fixed biased surrogate

The convergence analysis of Section A.2 takes the outer expectation in $\text{KL}(\mu\|\nu) + \lambda X_\mu(\mu\|\nu)$ under the exact target. To isolate surrogate bias, fix a smooth density $\tilde{\mu}$ that does not vary with ν during this analysis and define

$$\tilde{\mathcal{L}}[\nu] = \int_{\mathbb{R}^d} \tilde{\mu}(x) \log w(x) dx + \lambda \iint_{\mathbb{R}^d \times \mathbb{R}^d} \tilde{\mu}(x) \tilde{\mu}(y) |\log w(x) - \log w(y)| dx dy, \quad (46)$$

where the pointwise ratio $w(x) = \mu(x)/\nu(x)$ is still evaluated against the exact target. Repeating the selected-subgradient calculation of Section A.1 with μ replaced by $\tilde{\mu}$ only in the outer integrals gives

$$\text{grad}^{\text{FR}}[\tilde{\mathcal{L}}](x) = \nu(x) - \tilde{\mu}(x)(1 + 2\lambda \tilde{S}(x)), \quad \tilde{S}(x) := \mathbb{E}_{y \sim \tilde{\mu}}[\text{sgn}(w(x) - w(y))].$$

A stationary density ν^* of this selected flow therefore satisfies

$$\nu^*(x) = \tilde{\mu}(x) \left[1 + 2\lambda \int_{\mathbb{R}^d} \tilde{\mu}(y) \text{sgn}\left(\frac{\mu(x)}{\nu^*(x)} - \frac{\mu(y)}{\nu^*(y)}\right) dy \right]. \quad (47)$$

At $\lambda = 0$, the equation reduces to $\nu^* = \tilde{\mu}$ and the residual error is $\text{KL}(\mu\|\tilde{\mu})$. For $\lambda > 0$, the rank correction $1 + 2\lambda \tilde{S}^*$ raises ν^* relative to $\tilde{\mu}$ where μ/ν^* is comparatively high and lowers it where the ratio is low. For $\lambda < \frac{1}{2}$, the bracket in (47) is positive because $\tilde{S}^* \in [-1, 1]$. Antisymmetry of the sign also gives

$$\int_{\mathbb{R}^d} \tilde{\mu}(x) \tilde{S}^*(x) dx = \mathbb{E}_{x, y \sim \tilde{\mu}} \left[\text{sgn}\left(\frac{\mu(x)}{\nu^*(x)} - \frac{\mu(y)}{\nu^*(y)}\right) \right] = 0,$$

hence every stationary point $\nu^* = \tilde{\mu}(1 + 2\lambda \tilde{S}^*)$ is a probability density.

Theorem 2 Assume that μ is absolutely continuous with respect to $\tilde{\mu}$ and that $\lambda \in [0, \frac{1}{2}]$. Then every stationary density ν^* of the selected subgradient flow for $\tilde{\mathcal{L}}$ in (46) satisfies

$$\text{KL}(\mu\|\nu^*) \leq \text{KL}(\mu\|\tilde{\mu}) - \lambda \iint_{\mathbb{R}^d \times \mathbb{R}^d} \tilde{\mu}(x) \tilde{\mu}(y) \left| \frac{\mu(x)}{\nu^*(x)} - \frac{\mu(y)}{\nu^*(y)} \right| dx dy. \quad (48)$$

Proof. Substituting the stationary identity $\nu^* = \tilde{\mu}(1 + 2\lambda \tilde{S}^*)$ into the difference of forward KLs and using $\mu = \nu^* \cdot (\mu/\nu^*) = \tilde{\mu}(1 + 2\lambda \tilde{S}^*)(\mu/\nu^*)$,

$$\text{KL}(\mu\|\tilde{\mu}) - \text{KL}(\mu\|\nu^*) = \int_{\mathbb{R}^d} \mu(x) \log(1 + 2\lambda \tilde{S}^*(x)) dx = \int_{\mathbb{R}^d} \tilde{\mu}(x) \frac{\mu(x)}{\nu^*(x)} u(x) \log u(x) dx,$$

where $u(x) := 1 + 2\lambda \tilde{S}^*(x) \geq 1 - 2\lambda > 0$, since $\tilde{S}^* \in [-1, 1]$ (a $\tilde{\mu}$ -average of signs) and $\lambda < \frac{1}{2}$. The elementary inequality $u \log u \geq u - 1$ for $u > 0$ then yields

$$\text{KL}(\mu\|\tilde{\mu}) - \text{KL}(\mu\|\nu^*) \geq 2\lambda \int_{\mathbb{R}^d} \tilde{\mu}(x) \frac{\mu(x)}{\nu^*(x)} \tilde{S}^*(x) dx.$$

Applying the rank-correlation argument of Lemma 1 verbatim with μ replaced by $\tilde{\mu}$ in the outer measure,

$$\int_{\mathbb{R}^d} \tilde{\mu}(x) \frac{\mu(x)}{\nu^*(x)} \tilde{S}^*(x) dx = \frac{1}{2} \iint_{\mathbb{R}^d \times \mathbb{R}^d} \tilde{\mu}(x) \tilde{\mu}(y) \left| \frac{\mu(x)}{\nu^*(x)} - \frac{\mu(y)}{\nu^*(y)} \right| dx dy,$$

which gives (48). $\square$

At $\lambda = 0$, the bound is attained by $\nu^* = \tilde{\mu}$. For $0 < \lambda < \frac{1}{2}$, the subtracted discrepancy is nonnegative and is positive unless μ/ν^* is constant on the support of $\tilde{\mu}$. Thus, within the theorem's assumptions, X_μ improves the stationary-error bound for a fixed biased surrogate.

The threshold $\lambda < \frac{1}{2}$ is a sufficient condition for strict positivity of ν^* . The experiments use $\lambda = 1$ and do not extend the theorem beyond its stated range. The live surrogate also changes with the current flow and is represented by finite batches, so it is not the fixed smooth density assumed here.

References

- Tara Akhound-Sadegh, Jarrod Rector-Brooks, Avishek Joey Bose, Sarthak Mittal, Pablo Lemos, Cheng-Hao Liu, Marcin Sendera, Siamak Ravanbakhsh, Gauthier Gidel, Yoshua Bengio, Nikolay Malkin, and Alexander Tong. Iterated denoising energy matching for sampling from Boltzmann densities. In *International Conference on Machine Learning*, 2024.
- Tara Akhound-Sadegh, Jungyoon Lee, Avishek Joey Bose, Valentin De Bortoli, Arnaud Doucet, Michael Bronstein, Dominique Beaini, Siamak Ravanbakhsh, Kirill Neklyudov, and Alexander Tong. Progressive inference-time annealing of diffusion models for sampling from Boltzmann densities. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025.
- Michael Samuel Albergo and Eric Vanden-Eijnden. NETS: A non-equilibrium transport sampler. In *International Conference on Machine Learning*, pages 1026–1055. PMLR, 2025. arXiv preprint arXiv:2410.02711.
- Shun-ichi Amari. *Information Geometry and Its Applications*, volume 194 of *Applied Mathematical Sciences*. Springer, 2016.
- Luigi Ambrosio, Nicola Gigli, and Giuseppe Savaré. *Gradient Flows in Metric Spaces and in the Space of Probability Measures*. Lectures in Mathematics ETH Zürich. Birkhäuser, 2 edition, 2008.
- Michael Arbel, Alexander G. D. G. Matthews, and Arnaud Doucet. Annealed flow transport Monte Carlo. In *International Conference on Machine Learning (ICML)*, 2021.
- Dominique Bakry, Ivan Gentil, and Michel Ledoux. *Analysis and Geometry of Markov Diffusion Operators*, volume 348 of *Grundlehren der mathematischen Wissenschaften*. Springer, 2014.
- Martin Bauer, Martins Bruveris, and Peter W. Michor. Uniqueness of the Fisher–Rao metric on the space of smooth densities. *Bulletin of the London Mathematical Society*, 48(3):499–506, 2016.
- Julius Berner, Lorenz Richter, and Karen Ullrich. An optimal control perspective on diffusion-based generative modeling. *Transactions on Machine Learning Research*, 2024.
- David A. Case, H. Metin Aktulga, Kellon Belfon, David S. Cerutti, G. Andrés Cisneros, Vinícius W. D. Cruzeiro, Niel Forouzesh, Timothy J. Giese, Andreas W. Götz, Holger Gohlke, Saeed Izadi, Koushik Kasavajhala, Mehmet C. Kaymak, Edward King, Tom Kurtzman, Taisung S. Lee, Pengfei Li, Junmei Liu, Tyler Luchko, Ray Luo, Madushanka Manathunga, Matheus R. Machado, Huan M. Nguyen, Kurt A. O’Hearn, Alexey V. Onufriev, Fu Pan, Sergio Pantano, Rui Qi, Ali Rahnamoun, Daniel R. Roe, Adrian Roitberg, Celeste Sagui, Stephan Schott-Verdugo, Jian Shen, Carlos L. Simmerling, Nikolai R. Skrynnikov, Jeremy Smith, Jason Swails, Ross C. Walker, Junmei Wang, Haixin Wei, Romain M. Wolf, Xiongwu Wu, Yaxin Xue, Darrin M. York, Si Zhao, and Peter A. Kollman. AmberTools. *Journal of Chemical Information and Modeling*, 63(20):6183–6191, 2023. doi: 10.1021/acs.jcim.3c01153.
- Omar Chehab, Anna Korba, Austin Stromme, and Adrien Vacher. Provable convergence and limitations of geometric tempering for Langevin dynamics. In *International Conference on Learning Representations*, volume 2025, pages 5451–5493, 2025.
- Wendy D. Cornell, Piotr Cieplak, Christopher I. Bayly, Ian R. Gould, Kenneth M. Merz, David M. Ferguson, David C. Spellmeyer, Thomas Fox, James W. Caldwell, and Peter A. Kollman. A second generation force field for the simulation of proteins, nucleic acids, and organic molecules. *Journal of the American Chemical Society*, 117(19):5179–5197, 1995.
- Pierre Del Moral, Arnaud Doucet, and Ajay Jasra. Sequential Monte Carlo samplers. *Journal of the Royal Statistical Society: Series B*, 68(3):411–436, 2006.
- Arnaud Doucet, Nando De Freitas, and Neil Gordon. An introduction to sequential Monte Carlo methods. In *Sequential Monte Carlo methods in practice*, pages 3–14. Springer, 2001.
- Simon Duane, A. D. Kennedy, Brian J. Pendleton, and Duncan Roweth. Hybrid Monte Carlo. *Physics Letters B*, 195(2):216–222, 1987. doi: 10.1016/0370-2693(87)91197-X.

- Peter Eastman, Jason Swails, John D. Chodera, Robert T. McGibbon, Yutong Zhao, Kyle A. Beauchamp, Lee-Ping Wang, Andrew C. Simmonett, Matthew P. Harrigan, Chaya D. Stern, Rafal P. Wiewiora, Bernard R. Brooks, and Vijay S. Pande. OpenMM 7: rapid development of high performance algorithms for molecular dynamics. *PLoS Computational Biology*, 13(7):e1005659, 2017. doi: 10.1371/journal.pcbi.1005659.
- Víctor Elvira, Luca Martino, and Christian P Robert. Rethinking the effective sample size. *International Statistical Review*, 90(3):525–550, 2022.
- Loris Felardos, Jérôme Hénin, and Guillaume Charpiat. Designing losses for data-free training of normalizing flows on Boltzmann distributions. *arXiv preprint arXiv:2301.05475*, 2023.
- Daan Frenkel and Berend Smit. *Understanding Molecular Simulation: From Algorithms to Applications*, volume 1 of *Computational Science Series*. Academic Press, 2 edition, 2002.
- Marylou Gabrié, Grant M Rotskoff, and Eric Vanden-Eijnden. Adaptive Monte Carlo augmented with normalizing flows. *Proceedings of the National Academy of Sciences*, 119(10):e2109420119, 2022.
- Charles J. Geyer. Markov chain Monte Carlo maximum likelihood. In *Computing Science and Statistics: Proc. 23rd Symposium on the Interface*, pages 156–163, 1991.
- W. K. Hastings. Monte Carlo sampling methods using Markov chains and their applications. *Biometrika*, 57(1):97–109, 1970.
- Aaron Havens, Benjamin Kurt Miller, Bing Yan, Carles Domingo-Enrich, Anuroop Sriram, Brandon Wood, Daniel Levine, Bin Hu, Brandon Amos, Brandon Karrer, Xiang Fu, Guan-Horng Liu, and Ricky T. Q. Chen. Adjoint sampling: highly scalable diffusion samplers via adjoint matching. In *International Conference on Machine Learning (ICML)*, 2025.
- Jiajun He, Yuanqi Du, Francisco Vargas, Dinghuai Zhang, Shreyas Padhy, RuiKang OuYang, Carla Gomes, and José Miguel Hernández-Lobato. No trick, no treat: pursuits and challenges towards simulation-free training of neural samplers. In *ICLR 2025 Workshop on Frontiers in Probabilistic Inference*, 2025.
- William G. Hoover. Canonical dynamics: Equilibrium phase-space distributions. *Physical Review A*, 31(3):1695–1697, 1985. doi: 10.1103/PhysRevA.31.1695.
- Koji Hukushima and Koji Nemoto. Exchange Monte Carlo method and application to spin glass simulations. *Journal of the Physical Society of Japan*, 65(6):1604–1608, 1996.
- Araz Jakalian, David B. Jack, and Christopher I. Bayly. Fast, efficient generation of high-quality atomic charges. AM1-BCC model: II. parameterization and validation. *Journal of Computational Chemistry*, 23(16):1623–1641, 2002. doi: 10.1002/jcc.10128.
- Richard Jordan, David Kinderlehrer, and Felix Otto. The variational formulation of the Fokker–Planck equation. *SIAM Journal on Mathematical Analysis*, 29(1):1–17, 1998.
- Johannes Kästner. Umbrella sampling. *WIREs Computational Molecular Science*, 1(6):932–942, 2011. doi: 10.1002/wcms.66.
- Durk P Kingma and Prafulla Dhariwal. Glow: Generative flow with invertible 1x1 convolutions. *Advances in neural information processing systems*, 31, 2018.
- Leon Klein and Frank Noé. Transferable Boltzmann generators. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.
- Leon Klein, Andreas Krämer, and Frank Noé. Equivariant flow matching. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- Jonas Köhler, Leon Klein, and Frank Noé. Equivariant flows: exact likelihood generative learning for symmetric densities. In *International Conference on Machine Learning (ICML)*, 2020.

- David P. Landau and Kurt Binder. *A Guide to Monte Carlo Simulations in Statistical Physics*. Cambridge University Press, 4 edition, 2014. doi: 10.1017/CBO9781139696463.
- Paul Langevin. Sur la théorie du mouvement brownien. *Comptes rendus hebdomadaires des séances de l’Académie des sciences*, 146:530–533, 1908.
- Ben Leimkuhler and Charles Matthews. *Molecular Dynamics: With Deterministic and Stochastic Numerical Methods*, volume 39 of *Interdisciplinary Applied Mathematics*. Springer, 2015. doi: 10.1007/978-3-319-16375-8.
- Jun S Liu. *Monte Carlo strategies in scientific computing*, volume 10. Springer, 2001.
- Yulong Lu, Jianfeng Lu, and James Nolen. Accelerating Langevin sampling with birth–death. *arXiv preprint arXiv:1905.09863*, 2019.
- Alex Matthews, Michael Arbel, Danilo J. Rezende, and Arnaud Doucet. Continual repeated annealed flow transport Monte Carlo. In *ICML*, 2022.
- Nicholas Metropolis, Arianna W. Rosenbluth, Marshall N. Rosenbluth, Augusta H. Teller, and Edward Teller. Equation of state calculations by fast computing machines. *The Journal of Chemical Physics*, 21(6):1087–1092, 1953.
- Laurence Illing Midgley, Vincent Stimper, Gregor NC Simm, Bernhard Schölkopf, and José Miguel Hernández-Lobato. Flow annealed importance sampling bootstrap. In *NeurIPS 2022 AI for Science: Progress and Promises*, 2022.
- Muhammad Ferjad Naeem, Seong Joon Oh, Youngjung Uh, Yunjey Choi, and Jaesun Yoo. Reliable fidelity and diversity metrics for generative models. In *International conference on machine learning*, pages 7176–7185. PMLR, 2020.
- Christian A. Naesseth, Fredrik Lindsten, and David M. Blei. Markovian score climbing: variational inference with $\text{KL}(p||q)$. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
- Radford M. Neal. Annealed importance sampling. *Statistics and Computing*, 11(2):125–139, 2001.
- Frank Noé, Simon Olsson, Jonas Köhler, and Hao Wu. Boltzmann generators: Sampling equilibrium states of many-body systems with deep learning. *Science*, 365(6457):eaaw1147, 2019.
- Shūichi Nosé. A unified formulation of the constant temperature molecular dynamics methods. *The Journal of Chemical Physics*, 81(1):511–519, 1984. doi: 10.1063/1.447334.
- Nikolas Nüsken and Lorenz Richter. Solving high-dimensional Hamilton–Jacobi–Bellman PDEs using neural networks: perspectives from controlled diffusions and measures on path space. *Partial Differential Equations and Applications*, 2:48, 2021.
- Alexey Onufriev, Donald Bashford, and David A. Case. Exploring protein native states and large-scale conformational changes with a modified generalized Born model. *Proteins: Structure, Function, and Bioinformatics*, 55(2):383–394, 2004. doi: 10.1002/prot.20033.
- Felix Otto. The geometry of dissipative evolution equations: the porous medium equation. *Communications in Partial Differential Equations*, 26(1–2):101–174, 2001.
- Felix Otto and Cédric Villani. Generalization of an inequality by Talagrand and links with the logarithmic Sobolev inequality. *Journal of Functional Analysis*, 173(2):361–400, 2000.
- George Papamakarios, Theo Pavlakou, and Iain Murray. Masked autoregressive flow for density estimation. *Advances in neural information processing systems*, 30, 2017.
- Danyal Rehman, Oscar Davis, Jiarui Lu, Jian Tang, Michael Bronstein, Yoshua Bengio, Alexander Tong, and Avishek Joey Bose. Efficient regression-based training of normalizing flows for Boltzmann generators. In *International Conference on Learning Representations (ICLR)*, 2026.

- Danilo Rezende and Shakir Mohamed. Variational inference with normalizing flows. In *International conference on machine learning*, pages 1530–1538. PMLR, 2015.
- Lorenz Richter and Julius Berner. Improved sampling via learned diffusions. In *International Conference on Learning Representations (ICLR)*, 2024.
- Christian P. Robert and George Casella. *Monte Carlo Statistical Methods*. Springer Texts in Statistics. Springer, 2 edition, 2004. doi: 10.1007/978-1-4757-4145-2.
- Michael Schaefer and Martin Karplus. A comprehensive analytical treatment of continuum electrostatics. *The Journal of Physical Chemistry*, 100(5):1578–1599, 1996. doi: 10.1021/jp9521621.
- Henrik Schopmans and Pascal Friederich. Temperature-annealed Boltzmann generators. In *International Conference on Machine Learning (ICML)*, 2025. arXiv:2501.19077.
- Henrik Schopmans, Christopher von Klitzing, and Pascal Friederich. Efficient training of Boltzmann generators using off-policy log-dispersion regularization. *arXiv preprint arXiv:2602.03729*, 2026.
- John Skilling. Nested sampling for general Bayesian computation. *Bayesian Analysis*, 1(4):833–859, 2006.
- Nikola Surjanovic, Saifuddin Syed, Alexandre Bouchard-Côté, and Trevor Campbell. Parallel tempering with a variational reference. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
- Robert H. Swendsen and Jian-Sheng Wang. Replica Monte Carlo simulation of spin-glasses. *Physical Review Letters*, 57(21):2607–2609, 1986.
- Saifuddin Syed, Alexandre Bouchard-Côté, George Deligiannidis, and Arnaud Doucet. Non-reversible parallel tempering: a scalable highly parallel MCMC scheme. *Journal of the Royal Statistical Society Series B*, 84(2):321–350, 2022.
- Charlie B. Tan, Avishek Joey Bose, Chen Lin, Leon Klein, Michael M. Bronstein, and Alexander Tong. Scalable equilibrium sampling with sequential Boltzmann generators. In *International Conference on Machine Learning (ICML)*, 2025.
- Glenn M. Torrie and John P. Valleau. Nonphysical sampling distributions in Monte Carlo free-energy estimation: Umbrella sampling. *Journal of Computational Physics*, 23(2):187–199, 1977. doi: 10.1016/0021-9991(77)90121-8.
- Lorenz Vaitl, Kim A. Nicoli, Shinichi Nakajima, and Pan Kessel. Gradients should stay on path: better estimators of the reverse- and forward-KL divergence for normalizing flows. *Machine Learning: Science and Technology*, 2022.
- Francisco Vargas, Will Grathwohl, and Arnaud Doucet. Denoising diffusion samplers. In *International Conference on Learning Representations*, 2023.
- Cédric Villani. *Optimal Transport: Old and New*, volume 338 of *Grundlehren der mathematischen Wissenschaften*. Springer, 2009.
- Annika von Klitzing, Denis Blessing, Henrik Schopmans, Pascal Friederich, and Gerhard Neumann. Learning Boltzmann generators via constrained mass transport. In *International Conference on Learning Representations (ICLR)*, 2026.
- Martin J. Wainwright and Michael I. Jordan. Graphical models, exponential families, and variational inference. *Foundations and Trends in Machine Learning*, 1(1–2):1–305, 2008.
- Junmei Wang, Romain M. Wolf, James W. Caldwell, Peter A. Kollman, and David A. Case. Development and testing of a general Amber force field. *Journal of Computational Chemistry*, 25(9):1157–1174, 2004. doi: 10.1002/jcc.20035.

- Peter Wirnsberger, Andrew J. Ballard, George Papamakarios, Stuart Abercrombie, Sébastien Racanière, Alexander Pritzel, Danilo Jimenez Rezende, and Charles Blundell. Targeted free energy estimation via learned mappings. *The Journal of Chemical Physics*, 153(14), 2020.
- Dongyeop Woo and Sungsoo Ahn. Iterated energy-based flow matching for sampling from Boltzmann densities. In *NeurIPS*, 2024. arXiv:2408.16249.
- Dawn B. Woodard, Scott C. Schmidler, and Mark Huber. Conditions for rapid mixing of parallel and simulated tempering on multimodal distributions. *The Annals of Applied Probability*, 19(2):617–640, 2009.
- Hao Wu, Jonas Köhler, and Frank Noé. Stochastic normalizing flows. *Advances in neural information processing systems*, 33:5933–5944, 2020.
- Leo Zhang, Peter Potapchik, Jiajun He, Yuanqi Du, Arnaud Doucet, Francisco Vargas, Hai-Dang Dau, and Saifuddin Syed. Accelerated parallel tempering via neural transports. *arXiv preprint arXiv:2502.10328*, 2025.
- Qinsheng Zhang and Yongxin Chen. Path integral sampler: A stochastic control approach for sampling. In *International Conference on Learning Representations*, 2022.