

Mode Coverage in Boltzmann Generators via Log-Ratio Variation

Xuda Ye (Purdue)

Joint work with Qi Feng, Rongjie Lai and Di Qi

September 9, 2026

Mathematics of Dynamical Systems in Modern Machine Learning
Brin Mathematics Research Center, UMD


The Sampling Problem, and Where It Fails

- Given a potential U on $\mathbb{R}^d$, draw samples from

$$\pi(x) \propto e^{-U(x)},$$

with access to U and ∇U alone.

- U is multimodal: the mass splits among basins separated by barriers, and the quantity wanted is a ratio of masses between basins.
- Mode collapse**: the sampler occupies some basins and never reaches the rest. Nothing it reports shows the absence.



Boltzmann Generators: a Global Transport

One invertible map in place of a long chain. [Noé et al., *Science* 2019]

- Source $\pi_0 \propto e^{-U_0}$, Gaussian. Normalizing flow G trained so that $G_{\#}\pi \approx \pi_0$.
- Pushforward $\nu := (G^{-1})_{\#}\pi_0$ has a tractable density $\nu(y) = \pi_0(G(y))|\det J_G(y)|$.
- Weight $w(y) = \pi(y)/\nu(y)$, and the log-ratio $z = \log w$:

$$z(y) = U_0(G(y)) - U(y) - \log|\det J_G(y)|.$$

- Forward KL is exactly $\text{KL}(\pi\|\nu) = \mathbb{E}_{y\sim\pi}[z(y)]$.
- Reweighting by w corrects the estimate where the flow puts samples, **never a region it misses**. Effective sample size $\text{ESS}(w) = (\sum_i w_i)^2 / (N \sum_i w_i^2)$ is read on the support reached.

The Log-Ratio Variation

Wanted: a term added to the forward KL that evaluates at points from any distribution we can generate, needs **no normalizing constant** of π , and keeps $\nu = \pi$ a minimizer.

Log-ratio variation at a weighting distribution ω

$$X_{\omega}(\pi\|\nu) := \mathbb{E}_{y, y' \sim \omega} [|z(y) - z(y')|], \quad z = \log \frac{\pi}{\nu}.$$

Mean absolute pairwise difference of the log-ratio, both points drawn from ω .

- Normalizers cancel. A constant shift of z leaves every difference unchanged.
- $X_{\omega}(\pi\|\nu) = 0$ exactly when z is constant ω -a.s., i.e. $\nu \propto \pi$ on the support of ω .
- ω says where flatness is demanded, and it need not be π . This is the whole freedom the design uses.

The KLXX Loss

Two log-ratio variations appended to the forward KL, the two X's:

$$\begin{aligned}\text{KLXX}[G] &= \underbrace{\text{KL}(\pi\|\nu)}_{\text{target fit}} + \underbrace{X_{\pi}(\pi\|\nu)}_{\text{target variation}} + \underbrace{X_{(\hat{\pi}+\bar{\nu})/2}(\pi\|\nu)}_{\text{mixture variation}} \\ &= \mathbb{E}_{y\sim\pi}[z(y)] + \mathbb{E}_{y,y'\sim\pi}[|z(y) - z(y')|] + \mathbb{E}_{y,y'\sim(\hat{\pi}+\bar{\nu})/2}[|z(y) - z(y')|].\end{aligned}$$

- $\hat{\pi}$: the quench and temper (QT) distribution, supplying candidate modes.
- $\bar{\nu}$: the detached pushforward, equal to ν but outside the differentiation graph.

Generalized loss, mixture $\xi := \alpha\hat{\pi} + \beta\bar{\nu}$, with $\lambda, \vartheta, \alpha, \beta \geq 0$, $\alpha + \beta = 1$:



$$\mathcal{L}[G] = \text{KL}(\pi\|\nu) + \lambda X_{\pi}(\pi\|\nu) + \vartheta X_{\alpha\hat{\pi}+\beta\bar{\nu}}(\pi\|\nu).$$

KLXX is determined by $(\lambda, \vartheta, \alpha, \beta) = (1, 1, \frac{1}{2}, \frac{1}{2})$, written $\text{KL}+X_{\pi}+X_{(\hat{\pi}+\bar{\nu})/2}$.

Quench and Temper: Finding Candidate Modes


Three steps on each source sample $y \sim \pi_0$. Oracle access to U and ∇U only.

- 1 Melt: $y \leftarrow y + m_e \zeta$, $\zeta \sim \mathcal{N}(0, I_d)$, at melt scale $m_e > 0$.
- 2 Quench: $\dot{y} = -\nabla U(y)$, into whichever basin the melted particle landed in.
- 3 Temper: $\dot{y} = -\nabla U(y) + \sqrt{2} \dot{B}_t$, a spread around that basin.



Why Mix QT with the Pushforward

- $\hat{\pi}$ is inaccurate everywhere. Admissible: a variation asks for flatness of z , never for the mass at those points.
- $X_{\hat{\pi}}$ weights points inside the basins only, and says nothing where ν puts mass and $\hat{\pi}$ does not: the **intermodal** regions.



The Collision: Why Mixture Handles Collapse?

Under the **mixture** $\xi = \frac{1}{2}(\hat{\pi} + \bar{\nu})$, training drives

$$z(y) = \log \frac{\pi(y)}{\nu(y)} \approx \log \frac{\pi(y')}{\nu(y')} = z(y'), \quad y \sim \hat{\pi}, \quad y' \sim \bar{\nu}.$$

- y in a target mode: $z(y) \rightarrow +\infty$ if ν **misses** it.
- y' where the flow puts mass: $z(y') \rightarrow -\infty$ if ν **leaks**.
- Both hold only if ν covers every mode and avoids the intermodal regions.

KLXX Training Algorithm

Training of the generalized KLXX loss with parameters $(\lambda, \vartheta, \alpha, \beta)$.

Source samples $\mathcal{Y}_0 \sim \pi_0$ and QT samples $\hat{\mathcal{Y}} = \text{QT}(\mathcal{Y}_0, U)$, fixed before training.

Per gradient step, run:

- 1 pushforward batch $\mathcal{X}_\nu = G^{-1}(\mathcal{X}_0)$, from a random $\mathcal{X}_0 \subset \mathcal{Y}_0$ of size B ;
- 2 target surrogate $\mathcal{X}_\pi = \text{SMC}(\mathcal{X}_\nu)$, from pushforward ν to π ;
- 3 mixture batch $\mathcal{X}_\xi \sim \xi = \alpha\hat{\pi} + \beta\bar{\nu}$, drawn from $\hat{\mathcal{Y}}$ and $\mathcal{X}_\nu$.

$$\widehat{\mathcal{L}}[G] = \underbrace{\widehat{\mathbb{E}}_{y \in \mathcal{X}_\pi} [z(y)]}_{\text{forward KL}} + \underbrace{\lambda \widehat{\mathbb{E}}_{y, y' \in \mathcal{X}_\pi} [|z(y) - z(y')|]}_{\text{target variation}} + \underbrace{\vartheta \widehat{\mathbb{E}}_{y, y' \in \mathcal{X}_\xi} [|z(y) - z(y')|]}_{\text{mixture variation}}$$

One Adam step on $\widehat{\mathcal{L}}[G]$ updates G ; only z carries the gradient.

The Empirical Averages: Exact Calculation

- On a batch $\mathcal{X} = \{y_1, \dots, y_B\}$, the empirical averages are

$$\widehat{\mathbb{E}}_{y \in \mathcal{X}}[z(y)] := \frac{1}{B} \sum_{i=1}^B z(y_i),$$

$$\widehat{\mathbb{E}}_{y, y' \in \mathcal{X}}[|z(y) - z(y')|] := \frac{2}{B(B-1)} \sum_{1 \leq i < j \leq B} |z(y_i) - z(y_j)|.$$

- Each $z(y_i)$ is a scalar. Sorting $z_{(1)} \leq \dots \leq z_{(B)}$ gives the double sum exactly,

$$\sum_{i < j} |z(y_i) - z(y_j)| = \sum_{k=1}^B (2k - B - 1) z_{(k)},$$

at cost $O(B \log B)$: the $\binom{B}{2}$ pairs are never formed.

Adaptive Staging Boltzmann Generator: Interpolation

One flow from π_0 to π is not enough in high dimension. Interpolate the potentials and train one flow per [stage](#):

$$U_t = (1 - t)U_0 + tU, \quad \pi_t \propto e^{-U_t}, \quad 0 = t_0 < t_1 < \dots < t_K = 1.$$

- Stage k carries $\pi_{k-1} \rightarrow \pi_k$, a deformation small enough for one flow.
- The schedule cannot be fixed in advance: no target data, no knowledge of where the wells are. The difficulty is unknown until it is attempted.
- **Accept a stage only when its sample ESS clears a gate τ_v** , otherwise shrink the increment and retry.
- Easy interpolations take large stages, hard ones take small.

Adaptive Staging Boltzmann Generator: Algorithm

Sample set $\mathcal{Y}_0 \sim \pi_0$, $t_0 = 0$. At the stage $t_{k-1} \mapsto t_k$:

- 1 propose t_k by extrapolating the last accepted increment,

$$t_k = \min(1, t_{k-1} + \Gamma(t_{k-1} - t_{k-2}));$$

- 2 build the **QT set** $\hat{\mathcal{Y}}_k = \text{QT}(\mathcal{Y}_{k-1}, U_{t_k})$; (prebuilt before training)

- 3 train G_k on $\pi_{k-1} \rightarrow \pi_k$ by the KLXX training algorithm using $\hat{\mathcal{Y}}_k$;

- 4 weights $w_k = \exp(z_k)$ on $\mathcal{Y}_{k-1}$ pushed through G_k^{-1} .

If $\text{ESS}(w_k) < \tau_v$: shrink $t_k \leftarrow t_{k-1} + \gamma(t_k - t_{k-1})$ and repeat from 2.

Otherwise: resample, rejuvenate under U_{t_k} , and set $\mathcal{Y}_k$.

Defaults $\tau_v = 0.4$, $\Gamma = 1.5$, $\gamma = 0.7$. Output: the schedule $\{t_k\}$ and the stage maps $\{G_k\}$, which the generator replays on fresh source samples.

How the Error Propagates


Inference replays the accepted stages on N fresh source particles: push forward, reweight, resample, rejuvenate.

- For every bounded observable φ , with π_K^N the returned empirical measure, $w_k = \pi_k / \nu_k$ the stage weights and $\text{osc}(\varphi) = \text{ess sup } \varphi - \text{ess inf } \varphi$,

$$\|\pi_K^N \langle \varphi \rangle - \pi_K \langle \varphi \rangle\|_2 \leq \frac{\text{osc}(\varphi)}{2\sqrt{N}} \sum_{k=0}^K \prod_{j=k+1}^K 3\|w_j\|_\infty,$$


- **The scheme is asymptotically unbiased: its error is removed by raising N alone.** Classical Monte Carlo instead needs the simulation time to grow.

2D Benchmarks: Three-Well and Himmelblau



Three-Well (top), Himmelblau (bottom). Every method reports a high ESS.

2D Benchmarks: Where the ESS Is Earned





Coverage below 0.9: the ESS is earned only on the wells reached.

Product Multi-Well up to $d = 256$

$$U(\mathbf{x}) = \frac{1}{2}\|\mathbf{x}\|^2 + 12 \sum_{i=1}^k e^{-x_i^2}, \quad k = \lfloor \log_2 d \rfloor, \quad d = 16, \dots, 256.$$


- KL+L1: log-dispersion regularization penalty LDR-L1 (Schopmans et al.)
- FAB: flow annealed importance sampling bootstrap (Midgley et al.)



Every log-ratio variation beats forward KL at all sixteen d .

Coefficient Sweep at $d = 100$

$\text{KL} + \lambda X_\pi$ on the same potential, $\lambda = 0$ being forward KL.



KLXX (0.906, red line) sits above every value the sweep attains. Tuning one coefficient does not reproduce the mixture variation.

Boltzmann Generator on Three Molecules


- The difficulty is not multimodality alone: the Lennard–Jones core is **singular**, so an early proposal can land at energies far outside the thermally occupied region.
- Implicit solvent at 300 K, no target data, one seed. The path is regularized, ending short of the bare potential.

molecule (d)	KL+X $_{\pi}$		KLXX	
	$\hat{F}$	stages	$\hat{F}$	stages
<i>N</i> -methylacetamide, NMA (30)	2.06	6	1.77	5
glycerol (36)	3.77	7	3.28	6
neutral diethanolamine (48)	19.90	10	15.84	10

- Propagation factor $\hat{F} = \prod_k \text{ESS}_k^{-1/2}$, from the sample ESS of each accepted stage; smaller is better.

Dihedral Marginals against Parallel Tempering

Reference: parallel tempering, six replicas from 300 to 800 K. Densities symmetrized under $\varphi \mapsto -\varphi$, as the achiral targets allow.



NMA shows one well: its gauche population is too small to appear. Largest departure: glycerol trans, KLXX at about half of the $KL+X_\pi$ excess.

Conclusions

- Forward KL trains only on the samples it is given. When a mode is missing from the flow and from those samples, the ESS stays high and says nothing: a **fake ESS**.
- X_ω compares $\log(\pi/\nu)$ at pairs of points. It needs no normalizing constant, and it can be evaluated wherever we can draw samples: at π for accuracy, at $\hat{\pi}$ for the modes the flow has not found, at $\bar{\nu}$ to keep mass out of the gaps.
- The loss itself makes the flow **visit every mode**. The extra cost per step is two sorts and one more batch, plus the QT set built once.

- F. Noé, S. Olsson, J. Köhler, H. Wu. Boltzmann generators. *Science* 365 (2019).
- L. Midgley, V. Stimper, G. Simm, B. Schölkopf, J. Hernández-Lobato. Flow annealed importance sampling bootstrap. *ICLR* (2023).
- H. Schopmans, C. von Klitzing, P. Friederich. Efficient training of Boltzmann generators using off-policy log-dispersion regularization. arXiv:2602.03729 (2026).
- C. Durkan, A. Bekasov, I. Murray, G. Papamakarios. Neural spline flows. *NeurIPS* (2019). The architecture used at every stage here.